

PADMAKANYA MULTIPLE CAMPUS
BAGBAZAAR , KATHMANDU
INSTITUTE OF SCIENCE AND TECHNOLOGY



Final Year Project Report
On
Intellexa : A Content Analysis AI

Submitted to
Department of Computer Science and Information Technology
PadmaKanya Multiple Campus

In partial fulfillment of the requirements for Bachelor's Degree in Computer science and Information Technology

Submitted By:

Kopila Devkota(28371)

Sapana Khatiwada(28411)

November, 2025

LETTER OF RECOMMENDATION

It is with great pleasure that I recommend the project titled "Intellexa: A Content Analysis AI" for evaluation by Kopila Devkota (28371) and Sapana Khatiwada (28411). They are talented and dedicated students who have worked diligently on this project under my supervision. Their commitment, technical skills, and innovative approach have been truly commendable.

"Intellexa" reflects their ability to integrate advanced computational techniques, such as Retrieval-Augmented Generation (RAG) and AI-driven multi-format analysis, to address the practical challenge of information overload. The project demonstrates strong problem-solving and design capabilities in creating a unified platform for intelligent content interaction.

I am confident that this project will make a valuable contribution to the field and stands as a well-deserved accomplishment for these students. I wholeheartedly recommend that this project be processed for evaluation in partial fulfillment of the requirements for their Bachelor's degree of Science in Computer Science and Information Technology.

.....
Ramesh Singh Saud
Project Supervisor
Department of CSIT
PadmaKanya Multiple Campus

LETTER OF APPROVAL

This is to certify that the project report entitled "Intellexa: A Content Analysis AI" prepared by Kopila Devkota (28371) and Sapana Khatiwada (28411) in partial fulfillment of the requirements for the degree of Bachelor of Science in Computer Science and Information Technology has been thoroughly studied and executed. In our opinion, it is satisfactory in scope and quality and represents a significant achievement in their academic journey.

The project demonstrates the ability to design and implement an innovative AI-driven solution capable of processing and analyzing multiple content formats including videos, research papers, and web content through intelligent extraction and conversational interaction. It reflects strong technical proficiency, creativity, and problem-solving skills, meeting the high standards expected for the degree.

Evaluation Committee

.....
Sunita Shrestha
Program Coordinator
PadmaKanya Multiple Campus

.....
Ramesh Singh Saud
Project Supervisor
PadmaKanya Multiple Campus

.....
Internal Examiner
PadmaKanya Multiple Campus

.....
External Examiner
TU , IOST

DECLARATION

We hereby declare that the project report entitled "**Intellexa : A Content Analysis AI**" submitted to the **Department of Computer Science and Information Technology**, PadmaKanya Multiple Campus ,in partial fulfillment of the requirements for the degree of **Bachelor of Science in Computer Science and Information Technology (BSc. CSIT)**, is our own original work.

This work has been carried out under the supervision and guidance of **Mr. Ramesh Singh Saud**. To the best of our knowledge, the report contains no material previously published or written by another person, except where due reference is made in the text. All sources of information and knowledge used in this project have been duly acknowledged.

Submitted by:

Kopila Devkota (28371/078)

Sapana Khatiwada (28411/078)

Date:November2025

ACKNOWLEDGMENT

We would like to express our profound gratitude to all individuals who have supported and guided us throughout the journey of completing our project report titled "**Intellexa: A Content Analysis AI**". This accomplishment stands as a testament to the collective effort and encouragement we received from various quarters.

Our deepest appreciation goes to our project supervisor, **Mr. Ramesh Singh Saud**, whose unwavering guidance, invaluable insights, and constant motivation were instrumental in shaping this project. His expertise and mentorship provided us with the direction and confidence needed to navigate challenges and refine our work.

We are also grateful to the Head of CSIT department **Ms. Sunita Shrestha**, who impacted and inspired us. We are also sincerely thankful to the **Department of Computer Science and Information Technology at PadmaKanya Multiple Campus** for providing an enriching academic environment and the necessary resources to undertake this project. The foundational knowledge and technical skills imparted by the faculty have been crucial to the successful execution of this work.

Our heartfelt thanks extend to our friends and colleagues for their constructive feedback and moral support during the course of this project. Lastly, we owe a special debt of gratitude to our families for their unconditional love, patience, and encouragement, which kept us motivated and resilient throughout this endeavor.

This project is the culmination of collaborative spirit and shared wisdom, and we are eternally grateful to everyone who contributed to its realization, directly or indirectly.

ABSTRACT

In today's digital information landscape, learners and professionals face significant challenges in efficiently processing, analyzing, and extracting meaningful insights from diverse content formats including YouTube videos, research papers, PDF documents, and web articles. To address these challenges, **Intellexa: A Content Analysis AI** has been developed as an intelligent content analysis platform that enhances information comprehension through advanced AI-powered processing, multi-format support, and contextual understanding.

The system integrates several sophisticated components, including intelligent content routing using **LangGraph** state management, multi-format content processing for PDFs, research papers, and web articles through specialized web scraping and parsing modules, context-aware video analysis with temporal timestamping via **YouTube transcript processing**, and adaptive conversational AI powered by **Retrieval-Augmented Generation (RAG)** with semantic search capabilities. Together, these modules form a cohesive platform that helps users analyze, comprehend, and interact with diverse digital content more effectively.

Intellexa is implemented as a full-stack web application built with modern technologies including **Streamlit** for the responsive frontend interface, Python with LangGraph for the intelligent workflow engine, **Chroma vector database** for semantic search, and leverages advanced APIs like Groq's Llama-3.3-70b model for high-speed inference and **HuggingFace transformers** for embedding generation. The result is a scalable, interactive, and user-friendly system that empowers learners, researchers, and professionals to process information more intelligently, retain knowledge better, and achieve deeper understanding through AI-driven content analysis and multi-modal interaction capabilities.

Keywords: *LangGraph, YouTube transcript processing, Retrieval-Augmented Generation, Streamlit, Chroma vector database, HuggingFace transformers*

CONTENTS

LETTER OF RECOMMENDATION.....	2
LETTER OF APPROVAL	3
DECLARATION.....	4
ACKNOWLEDGMENT.....	5
ABSTRACT.....	6
TABLE OF FIGURES	9
LIST OF TABLES	9
LIST OF ABBREVIATIONS.....	11
CHAPTER 1: INTRODUCTION	12
1.1 Introduction.....	12
1.2 Problem Statement	13
1.3 Objectives	13
1.4 Scope and Limitations.....	14
1.4.1 Scope.....	14
1.4.2 Limitations	14
1.5 Development Methodology	14
1.6 Report Organization	15
Chapter 1 : Introduction	16
Chapter 2: Background Study and Literature Review	16
Chapter 3: System Analysis	16
Chapter 4: System Design.....	16
Chapter 5: Implementation and Testing	16
Chapter 6: Conclusion and Future Recommendations.....	16
CHAPTER 2: BACKGROUND STUDY AND LITERATURE REVIEW.....	17
2.1 Background Study.....	17
2.2 Literature Review.....	18
2.2.1 Content Analysis and Retrieval-Augmented Generation Systems.....	18
2.2.2 Intelligent Web Content Extraction and Analysis	18
2.2.3 Conversational AI and Educational Dialogue Systems	19
2.2.4 Audio Generation and Multimodal Interaction	19
2.2.5 Study of Existing Systems	19
CHAPTER 3: SYSTEM ANALYSIS	21

3.1 System Analysis	21
3.1.1 Requirement Analysis	21
3.1.2 Feasibility Analysis.....	24
3.1.3 Analysis.....	26
CHAPTER 4: SYSTEM DESIGN	30
4.1 Design Overview	30
4.1.1 System Architecture Diagram	30
4.1.2 Refined Class Diagram:	31
4.1.3 Refined Sequence Diagram.....	33
4.1.4 Refined Activity Diagram:	34
4.1.5 Deployment Diagram.....	36
4.2 Algorithm Details.....	36
4.2.1 Semantic Search with MMR (Maximal Marginal Relevance).....	36
4.2.2 LangGraph State Machine for Query Routing.....	38
4.2.3 Content-Type Specific Processing Pipeline	40
4.2.4 Edge TTS Audio Synthesis with Text Preprocessing.....	41
CHAPTER 5: IMPLEMENTATION AND TESTING	44
5.1 Implementation	44
5.1.1 Tools Used.....	44
5.1.2 Implementation Details of Modules.....	45
5.2 Testing	50
5.2.1 Test Cases for Unit Testing	50
5.2.2 Test Cases for System Testing.....	53
CHAPTER 6: CONCLUSION AND FUTURE RECOMMENDATION	55
6.1 : Conclusion	55
6.2 Recommendation	55
REFERENCES	57
APPENDIX.....	58

TABLE OF FIGURES

Figure 1: Prototyping Model.....	15
Figure 3. 1: Use Case Diagram	23
Figure 3. 2: Gantt Chart	26
Figure 3. 3: Class Diagram	27
Figure 3. 4: Sequence Diagram.....	28
Figure 3. 5: Activity Diagram	29
Figure 4. 1 : System Architecture Design	31
Figure 4. 2: Refined Class Diagram.....	32
Figure 4. 3: Refined Sequence Diagram	34
Figure 4. 4: Refined Activity Diagram	35
Figure 4. 5: Deployment Diagram	36
Figure 6. 1 : GitHub Repository Structure of Intellexa Project	58
Figure 6. 2:Dashboard Demonstrating Input and Load Youtube Video URL.....	58
Figure 6. 3:Providing prompt and getting response with Timestamp	59
Figure 6. 4: Load Web Url and getting web contents	59
Figure 6. 5: MultiLinguistic Functionality Demonstration.....	60
Figure 6. 6: Asking General query without loading URLs	60

LIST OF TABLES

Table 4. 1: Description Table For Refined Class Diagram	32
Table 5. 1: Core AI Technologies.....	44
Table 5. 2: Content Processing Libraries	45
Table 5. 3: AI / ML Integration	45
Table 5. 4: Unit Testing For Content Processing	51
Table 5. 5: Unit Testing for AI Engine.....	51
Table 5. 6: Unit Testing For Vector Store	52
Table 5. 7: Unit Testing For Audio Service.....	53
Table 5. 8: System Testing For Content Analysis AI System	54

LIST OF ABBREVIATIONS

AI	Artificial Intelligence
API	Application Programming Interface
BERT	Bidirectional Encoder Representations from Transformers.
CLIP	Contrastive Language-Image Pre-training
CNN	Convolutional Neural Network
DB	Database
LLM	Large Language Model
ML	Machine Learning
MMR	Maximal Marginal Relevance
NLP	Natural Language Processing
PDF	Portable Document Format
RAG	Retrieval-Augmented Generation
SDLC	Software Development Life Cycle
TTS	Text to Speech
ViT	Vision Transformer

CHAPTER 1: INTRODUCTION

1.1 Introduction

The project *Intellexa: A Content Analysis AI* represents an innovative approach to addressing the challenges of multi-format digital content processing. This intelligent system enables seamless analysis and interaction with diverse content types including YouTube videos, web articles, research papers, and PDF documents through specialized processing pipelines. Each content format is handled by dedicated modules :YouTube videos undergo transcript extraction with temporal analysis, web content is processed through advanced scraping techniques using Trafilatura and BeautifulSoup, while research papers are managed through arXiv integration and sophisticated PDF parsing.

At the core of Intellexa's functionality lies its advanced AI infrastructure, which integrates LangGraph for intelligent workflow orchestration, Groq for high-speed large language model inference, and ChromaDB for semantic vector storage and retrieval. The system features a sophisticated conversational interface that supports both text and voice interactions, allowing users to ask contextual questions, request summaries, and explore content through natural dialogue. The platform maintains session-aware conversations using LangGraph checkpoints and provides multilingual support for content processing in English, Hindi, French, Spanish, and German, making it accessible to a diverse global user base.

The application provides a streamlined user experience through a unified web interface built with Streamlit, where users can input content URLs, select processing options, and engage in interactive conversations with their materials. When users load content, the system automatically detects the format, processes it through the appropriate pipeline, and enables rich contextual questioning with features like timestamp-specific video references and document page citations. This integrated approach transforms passive content consumption into active, insight-driven exploration, supporting enhanced learning outcomes, research efficiency, and knowledge discovery across academic, professional, and personal use cases.

1.2 Problem Statement

In today's digital ecosystem, users face significant challenges in efficiently processing and extracting meaningful insights from the vast and diverse range of content available across multiple platforms. The exponential growth of digital materials including videos, research papers, PDF documents, and web articles has led to information overload, making it difficult for individuals to identify relevant information without investing considerable time and effort. Compounding this issue is the fragmentation of content across different formats, each requiring specialized tools and workflows. This lack of integration forces users to constantly switch between applications, increasing cognitive load, reducing productivity, and hindering a cohesive understanding of multi-format materials.

Furthermore, traditional content consumption methods offer limited interactivity and contextual awareness, restricting users from engaging dynamically with the material. Existing tools rarely support natural language interactions, such as asking clarifying questions, requesting summaries, or exploring content through conversational interfaces. The absence of features like timestamp-aware video analysis and cross-content contextual understanding prevents users from forming meaningful connections between different sources. These limitations highlight the urgent need for an intelligent, unified content analysis system that streamlines information processing, supports seamless multi-format integration, and enables interactive, context-aware engagement with digital content.

1.3 Objectives

The specific objectives of this project are as follows:

- To develop an intelligent multi-source content processor capable of automatically detecting, extracting and analyzing diverse formats (videos, PDFs, research papers, web content) with structured semantic understanding.
- To implement a context-aware conversational AI that maintains content-grounded dialogues with multi-modal interactions (text/voice) and real-time audio responses for enhanced user engagement.
- To design a scalable knowledge architecture integrating multiple AI technologies with efficient document processing, providing foundation for continuous expansion and additional data source integration.

1.4 Scope and Limitations

1.4.1 Scope

The project encompasses the development of a comprehensive AI system from initial analysis and design through final testing phases, enabling users to efficiently extract and analyze insights from diverse digital content sources. The system processes multiple content format including YouTube videos, research papers, PDF documents, and web articles through specialized extraction pipelines with advanced semantic understanding for accurate information retrieval.

It facilitates intelligent content interaction through timestamp-aware video analysis and conversational AI, supporting both text and voice inputs with contextual awareness across user sessions. By integrating multi-format content loading, analysis, and multilingual interaction within a unified platform, the system continuously enhances its processing capabilities to adapt to evolving content patterns and user engagement behaviors.

1.4.2 Limitations

Despite its advanced capabilities, the current version of Intellexa has certain limitations:

- The system's functionality depends on third-party APIs including Groq LLM and YouTube Transcript API, making it vulnerable to service disruptions, rate limits, or unexpected changes from external providers.
- High computational demands for AI features can slow response times during peak usage, and the interface optimized for desktop offers limited mobile responsiveness, affecting accessibility on smaller devices.
- Complex PDF layouts, websites with anti-scraping protections, or password-protected materials may prevent accurate content processing and analysis.

1.5 Development Methodology

The Intellexa project was developed using an **Iterative Prototyping** methodology, essential for integrating complex, non-deterministic AI components. The process was built on a **modular three-layer architecture**, a Streamlit frontend, a LangGraph backend for stateful orchestration, and specialized data processors allowing components to be developed and tested in isolation.

Core development involved continuous refinement of the **AI workflow engine**. We began with a simple chat prototype and iteratively evolved it into a sophisticated, graph-based

state machine. This included implementing an enhanced router that combines regex pattern matching for greetings with a structured LLM chain for complex intent classification, directing queries to specialized handlers for video Q&A, web content, or general conversation.

Parallel iterations focused on the **Data and Retrieval Pipeline**. The system started with basic YouTube transcript processing and was progressively enhanced into a multi-source RAG pipeline. This incorporated the `WebContentProcessor` for PDFs, arXiv papers, and news articles, and employed MMR (Max Marginal Relevance) retrieval to ensure diverse and relevant context for the LLM.

New features like **Conditional Audio Generation** were integrated using feature flags, allowing the TTS system to be developed and tested without disrupting core functionality. This iterative, cycle-based approach enabled continuous validation of each AI component, facilitating organic evolution from a minimal viable prototype into a fully-featured, multi-modal content analysis platform.

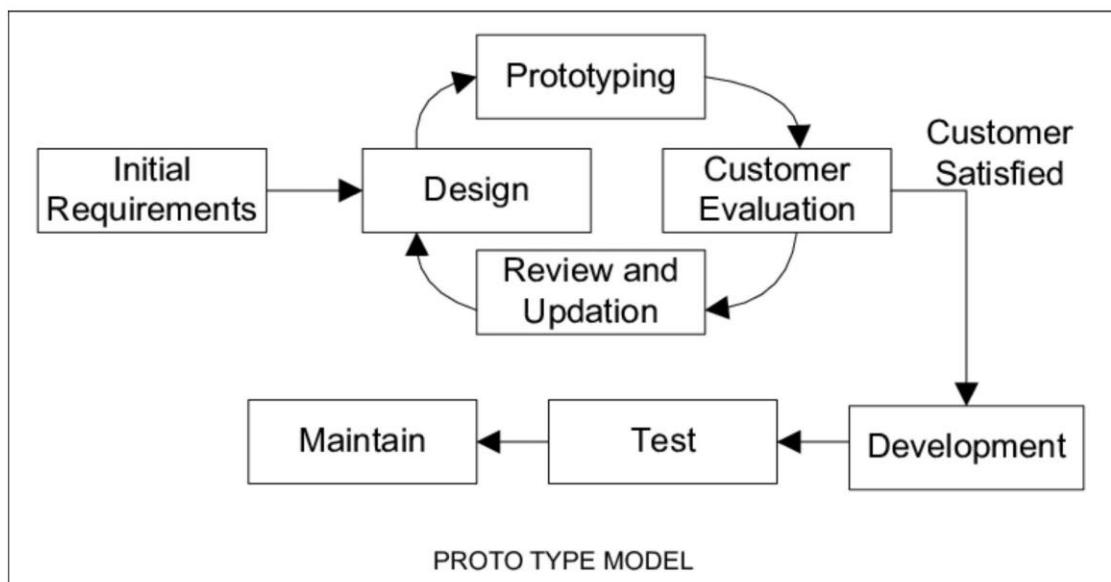


Figure 1: Prototyping Model

1.6 Report Organization

This report is organized into six chapters, each focusing on a specific aspect of the Intellexa Content Analysis AI System development:

Chapter 1 : Introduction

Introduces the project concept, problem statement in multi-format content analysis, objectives for intelligent content processing, scope and limitations of the system, and the prototyping methodology used for development.

Chapter 2: Background Study and Literature Review

Presents the background of AI in content analysis and education, studies of existing content processing systems, and reviews of related technologies including natural language processing, vector databases, and conversational AI that influenced the system design.

Chapter 3: System Analysis

Includes comprehensive requirement analysis for multi-format content processing, feasibility study covering technical, operational, and economic aspects, and overall system analysis to understand the challenges of integrating diverse content types before development.

Chapter 4: System Design

Describes the architecture of the proposed content analysis system, use-case diagrams for different user interactions, component structure and algorithms supporting core functionalities.

Chapter 5: Implementation and Testing

Provides detailed documentation of tools and technologies used in the stack, the system implementation process and unit and system test cases for content processing

Chapter 6: Conclusion and Future Recommendations

Summarizes the completed work in building a unified content analysis platform, key findings from user testing and performance evaluation, and suggests future improvements including expanded format support, enhanced multilingual capabilities, and advanced analytical features.

References and Appendix:

Lists all external sources referred to during the project and includes supporting materials such as diagrams, screenshots, and additional documentation.

CHAPTER 2: BACKGROUND STUDY AND LITERATURE REVIEW

2.1 Background Study

Artificial Intelligence (AI) has emerged as a transformative force in modern information processing, fundamentally reshaping how users interact with and extract value from diverse digital content. Through the integration of Machine Learning (ML), Natural Language Processing (NLP), and semantic analysis, AI enables intelligent content processing, adaptive information retrieval, and contextual understanding. These technologies collectively address the limitations of traditional content consumption by creating intelligent, user-centric analysis environments that enhance efficiency, information retention, and analytical productivity. AI-powered systems are capable of processing varied information sources from video transcripts and research papers to web documents to identify key insights and reduce cognitive load for users.

A fundamental concept underpinning this domain is multi-modal processing, where algorithms dynamically adapt their analytical techniques based on content type, format, and user intent. Modern systems employ a range of models, from transformer-based language understanding to specialized document parsing frameworks, to optimize content extraction. Research in educational psychology demonstrates that information is better understood and retained when analyzed in meaningful contexts, a principle known as the "processing depth effect." Contemporary AI systems build upon this by implementing semantic search, contextual embeddings, and Retrieval-Augmented Generation (RAG) to deliver relevant, context-aware information precisely when needed.

Complementing these capabilities are intelligent routing systems that guide users through complex content landscapes. Using hybrid approaches that combine rule-based classification, semantic analysis, and machine learning, these systems automatically detect content types, extract key information, and direct queries to specialized processing modules. Intellexa synthesizes these technological and cognitive foundations into a unified platform. The system processes multiple content formats through specialized parsers, implements intelligent query routing using LangGraph, facilitates contextual interactions via RAG-enhanced NLP, and provides multi-modal interfaces. By integrating these elements, Intellexa transforms challenges like content fragmentation and analysis inefficiency into structured, intelligent, and adaptive content analysis experiences.

2.2 Literature Review

2.2.1 Content Analysis and Retrieval-Augmented Generation Systems

The field of AI-powered content analysis has evolved significantly with the advent of advanced language models and retrieval systems. Intellexa builds upon several key technological approaches that have demonstrated effectiveness in processing and analyzing diverse content formats including videos, research papers, and web documents.

Retrieval-Augmented Generation (RAG) Systems

Retrieval-Augmented Generation represents a paradigm shift in how AI systems handle knowledge-intensive tasks by combining the strengths of dense retrieval with generative language models. Lewis et al. introduced RAG as a method that conditions on documents retrieved from a large corpus to generate more accurate and verifiable responses [1]. This approach has been particularly effective in educational and research contexts where factual accuracy and source verification are crucial. Intellexa implements RAG principles through its Chroma vector stores and HuggingFace embeddings, enabling context-aware responses grounded in specific video transcripts and web documents.

Multimodal Content Processing

Modern AI systems have increasingly focused on processing multiple content modalities. YouTube transcript analysis systems similar to Intellexa's implementation have been explored in educational contexts. Research by Kim et al. demonstrated that AI-powered video content analysis can significantly enhance learning outcomes by providing personalized explanations and timestamp-based references [2]. Intellexa extends this capability by processing not only video content but also diverse web resources including research papers, news articles, and PDF documents.

2.2.2 Intelligent Web Content Extraction and Analysis

Automated Document Processing

The challenge of extracting meaningful content from diverse web sources has been addressed through various computational approaches. The `trafilatura` library, which Intellexa employs for web content extraction, builds upon research by Barbaresi that demonstrated superior performance in main content extraction compared to traditional methods [3]. Similarly, `newspaper3k` implements algorithms for news-specific content extraction that have been validated in large-scale content analysis studies.

Research Paper Analysis Systems

ArXiv content processing, as implemented in Intellexa, follows established practices in academic information retrieval. Systems like Semantic Scholar and ArXiv-sanity have demonstrated the value of AI-powered paper recommendation and analysis [4]. Intellexa's approach to structuring research paper content with metadata alignment reflects best practices identified in systematic literature reviews of academic information systems.

2.2.3 Conversational AI and Educational Dialogue Systems

Language Model Applications in Education

The integration of large language models in educational systems has shown promising results. Recent studies by Kasneci et al. comprehensively analyzed the potential of ChatGPT and similar models in educational contexts, highlighting their capabilities in explanation generation, content summarization, and adaptive tutoring [5]. Intellexa leverages the Groq API with Llama-3.3-70b model to provide similar educational dialogue capabilities while maintaining context from specific content sources.

State Management in Conversational Systems

Intellexa's use of LangGraph for state management and workflow orchestration aligns with recent advancements in conversational AI architecture. Research by Zhou et al. demonstrated that structured state management significantly improves the coherence and context retention in extended educational dialogues [6]. The conditional routing and memory persistence implemented in Intellexa reflect these architectural best practices.

2.2.4 Audio Generation and Multimodal Interaction

Text-to-Speech in Educational Contexts

The integration of edge-tts for audio response generation in Intellexa builds upon research showing that multimodal presentation enhances learning retention. Studies by Mayer on multimedia learning principles confirm that combining textual and auditory information can improve knowledge transfer, particularly for complex concepts [7]. Intellexa's conditional audio generation provides this multimodal capability while maintaining user preference flexibility.

2.2.5 Study of Existing Systems

Commercial Content Analysis Platforms

Several commercial platforms offer content analysis capabilities that share similarities with Intellexa's features:

ChatPDF and Similar Tools: Systems like ChatPDF demonstrate the commercial viability of document-based Q&A systems [8]. These platforms allow users to interact with PDF documents through natural language queries. However, they typically focus exclusively on PDF processing and lack Intellexa's multimodal capabilities integrating video analysis, web content processing, and comprehensive content management.

YouTube Summary with ChatGPT: Various browser extensions and tools that combine YouTube transcript extraction with ChatGPT interfaces have gained popularity for video content analysis. While effective for basic summarization, these tools lack the sophisticated retrieval mechanisms, content persistence, and integrated workflow management that Intellexa provides through its LangGraph architecture.

ResearchRabbit and Connected Papers: These academic discovery tools demonstrate effective paper recommendation and analysis capabilities [9]. ResearchRabbit specifically visualizes research paper networks and provides content-based recommendations. Intellexa incorporates similar academic paper processing capabilities while extending them to general web content and video analysis.

Educational AI Assistants

Khanmigo (Khan Academy AI): Khan Academy's AI assistant represents a significant advancement in educational AI, providing personalized tutoring and explanations [10]. While Khanmigo excels in structured curriculum contexts, Intellexa differentiates itself by supporting user-provided content across diverse formats and sources, making it more adaptable to individual research needs and self-directed learning.

Consensus AI: This platform specializes in research paper analysis and evidence-based answers [11]. Consensus uses similar RAG approaches for scientific literature but focuses exclusively on research papers. Intellexa provides broader content analysis capabilities while maintaining research-grade accuracy for academic content.

In summary, while individual components of Intellexa have precedents in existing systems, its comprehensive integration of video analysis, web content processing, research paper understanding, and conversational AI within a unified architecture represents a significant advancement in content analysis systems. The platform's ability to maintain context across diverse content types while providing source-grounded responses addresses key limitations observed in more specialized tools.

CHAPTER 3: SYSTEM ANALYSIS

3.1 System Analysis

System analysis is a critical phase in the software development lifecycle, involving an in-depth study of the problem domain, identification of requirements, and assessment of the proposed solution's feasibility. The primary objective of this phase is to comprehensively understand the system's objectives and define the most efficient and user-friendly approach to achieve them.

For Intellexa, the system analysis phase involved examining the content processing challenges faced by students, researchers, and professionals, evaluating the limitations of existing tools, and designing an AI-powered solution that enhances content comprehension and analytical efficiency. This phase encompassed detailed requirement analysis, feasibility assessment, and comprehensive system evaluation to ensure the solution's practicality, cost-effectiveness, and technical viability.

The Intellexa platform was analyzed with emphasis on:

- Providing unified multi-format content processing (videos, PDFs, research papers, web content)
- Delivering contextual understanding with temporal and semantic awareness
- Enabling natural language interactions through conversational AI
- Supporting multi-modal interfaces including voice and text
- Ensuring scalable architecture with robust AI integration

The subsequent sections detail the system's functional and non-functional requirements, evaluate the solution's feasibility, and examine the strengths and limitations of the proposed system architecture.

3.1.1 Requirement Analysis

For Intellexa, the primary user base comprises students, researchers, academics, and professionals who require intelligent processing, analysis, and interaction with diverse digital content formats through an AI-powered platform. Requirement gathering followed agile, user-centered design principles with iterative validation cycles.

i. Functional Requirements

Functional requirements describe the behavior of the overall functionality of system.

FR1: Dashboard Management

FR1.1: The system shall display a main dashboard to the user upon running the application.

FR1.2: The dashboard shall display the real-time conversation chat history.

FR1.3: The dashboard shall provide clear visual indicators for system status (e.g., "Content Loaded," "Processing," "Ready").

FR2: Content Loading and Configuration

FR2.1: The system shall allow the user to select a content type (e.g., Video, PDF, Web Article) from a predefined list.

FR2.2: The system shall provide an input field for the user to specify a URL to the content.

FR2.3: The system shall allow the user to select a preferred language for the analysis and interaction.

FR2.4: The system shall initiate the content loading process upon user command (e.g., a "Load" button).

FR3: Content Processing

FR3.1: The system shall automatically detect the type of content provided to validate the user's selection.

FR3.2: The system shall extract textual data from the provided URL (e.g., generate a transcript for a video, scrape text from a PDF or webpage).

FR3.3: The system shall process the extracted text to generate semantic embeddings for deep contextual understanding.

FR4: Multi-Modal Interaction

FR4.1: The system shall allow the user to input queries via text / audio.

FR4.2: The system shall convert user audio queries into text for processing.

FR4.3: The system shall generate and display text-based responses to user queries.

FR4.4: The system shall convert text-based responses into audio output upon user request.

FR5: Conversation Management

FR5.1: The system shall maintain the context of the conversation within a single session.

FR5.2: The system shall provide a "Clear" or "Reset" function to allow the user to erase the current chat history and start a new session.

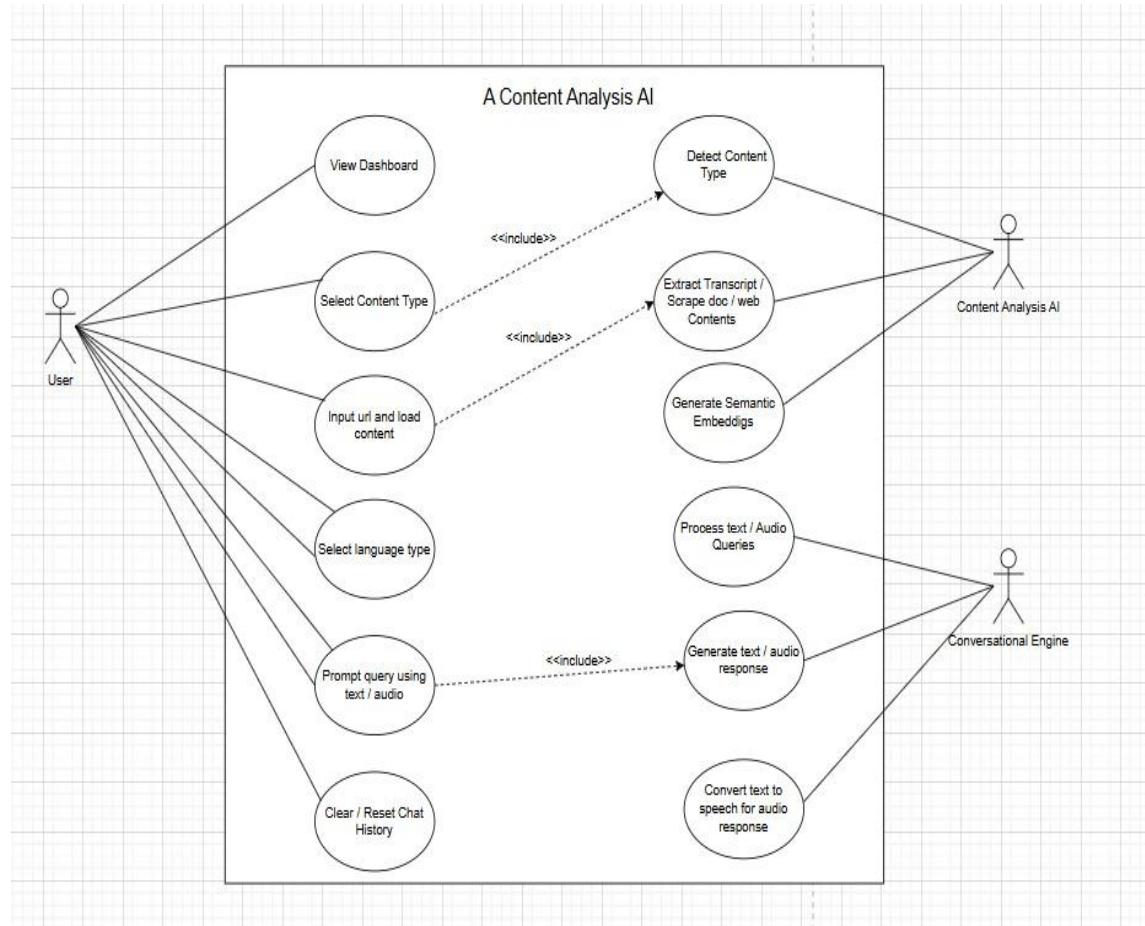


Figure 3. 1: Use Case Diagram

Based on the use case diagram, this Content Analysis AI system enables users to interact with various digital content through a conversational interface. Users begin by loading content such as YouTube videos, documents, or web pages via a dashboard, after which the system automatically processes it by detecting the content type, extracting and transcribing text, and generating semantic embeddings. The core functionality allows users to ask questions through text or voice prompts, leveraging a conversational engine to retrieve relevant information and provide answers in both text and synthesized audio, effectively bridging the gap between complex multimedia content and natural language interaction.

ii. Non Functional Requirements

NFR1: Performance & Scalability:

The system shall process content analysis requests with sub-5 second response times, support concurrent user sessions with consistent performance and maintain efficient memory utilization during content processing.

NFR2: Reliability & Availability:

The system shall maintain 99% uptime during operational periods, implement graceful degradation during external API failures, provide consistent performance across content types and maintain data integrity throughout processing pipelines.

NFR3 : Scalability & Maintainability:

The system shall support horizontal scaling to accommodate increasing user loads, implement modular architecture for seamless feature integration, provide comprehensive logging mechanisms for debugging and monitoring, and maintain clear separation between core processing and interface components.

NFR4: Usability & Accessibility:

The system shall provide an intuitive interface for multi-format content management, support accessibility standards for diverse user needs, offer clear navigation with responsive design, and deliver comprehensive user guidance with real-time feedback mechanisms.

3.1.2 Feasibility Analysis

Feasibility analysis comprises of determining if the proposed system is feasible, cost-effective, and viable with the resources and time constraints available. The Intellexa feasibility study involves the evaluation of the project according to technical, operational, and economic factors to ascertain if the idea is viable and can be carried out effectively.

i. Technical Feasibility

This evaluates whether the required technology and infrastructure are available to build the system.

- Technology Stack:** Streamlit framework provides responsive web interface with built-in UI components and real-time updates. Python with LangGraph enables sophisticated workflow management and stateful conversation handling as

backend. Chroma DB supports efficient semantic search and embedding storage for multi-format content and HuggingFace transformers are used as AI/ML Integration.

- **AI Integration:** Groq API (Llama-3.3-70b) provides high-speed inference for real-time content analysis. YouTube Transcript API enables temporal-aware video content processing and Multiple content processors (Trafilatura, BeautifulSoup, Newspaper3k) ensure robust web content extraction.
- **Content Processing Capabilities:** Advanced PDF parsing through PyPDF2 with hierarchical text extraction, Research paper analysis via specialized arXiv, integration and Multi-format document processing with metadata preservation.
- **Feasibility Verdict:** Technically feasible with the current development tools and available APIs.

ii. Operational Feasibility:

The Intellexa system is operationally feasible, offering an intuitive interface that allows users to effortlessly analyze diverse content from YouTube videos to research papers. By simply providing a URL, users can engage in natural language conversations, using text or voice, to get instant summaries and answers. The backend seamlessly automates complex processes like transcript extraction, semantic search, and audio synthesis. Its robust design reliably handles various content formats and includes clear error handling. This makes advanced AI-powered content analysis accessible and practical for non-technical users, significantly enhancing their information processing capabilities.

iii. Economic Feasibility

The Intellexa project is designed for high economic feasibility, requiring minimal financial investment as its core technologies are open-source. The implementation leverages cost-free libraries and frameworks like LangChain, LangGraph, and Streamlit, while utilizing ChromaDB for vector storage without licensing fees. By employing APIs like Groq for high-speed, low-cost inference and YouTube Transcript API for free data extraction, the system avoids expensive data acquisition and computational overhead. This strategic use of accessible tools and affordable cloud-based services ensures the development and maintenance costs remain low, making the project financially sustainable and viable for deployment.

iv. Schedule Feasibility

Schedule feasibility evaluates whether the system can be developed and deployed within the available timeframe, considering project complexity and resource availability.

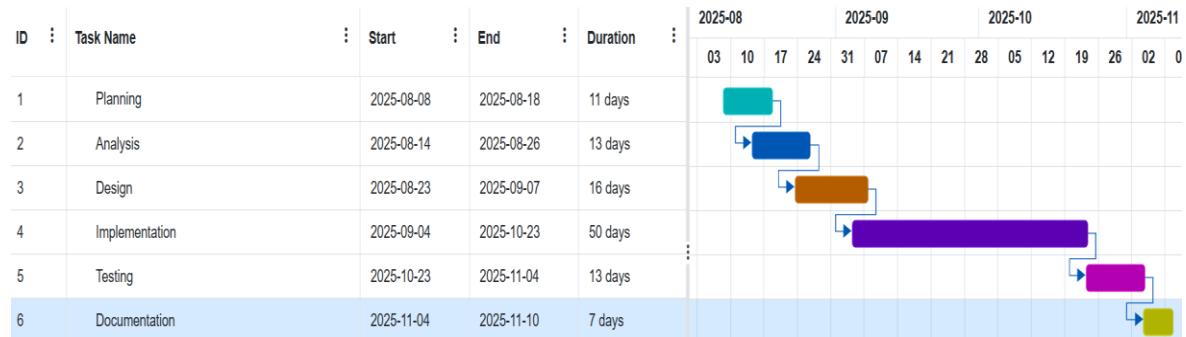


Figure 3. 2: Gantt Chart

To ensure the Intellexa project is developed in an efficient and timely fashion, a project timeline has been created. The timeline is divided into phases, and each phase takes a specific number of days. The Gantt chart above indicates the primary development phases and their estimated durations.

3.1.3 Analysis

In the Analysis phase, an object-oriented approach was employed to model the system architecture.

i. Object Modeling Using Class Diagram

Class diagram outlines the core components of the Intellexa Content Analysis AI System, showing how each class contributes to the overall functionality. The ContentProcessor handles multi-format content ingestion, while the VectorStore manages semantic storage and retrieval operations. The AIEngine orchestrates intelligent query processing and response generation, working alongside the User class which manages session interactions and preferences. The Message class maintains conversation history and session state, while the AudioService provides text-to-speech capabilities for multi-modal responses. Together, these classes form an integrated system that processes diverse content types, manages user interactions, and delivers intelligent responses through a structured, modular architecture.

The class diagram of Intellexa to demonstrate object modelling is as follows:

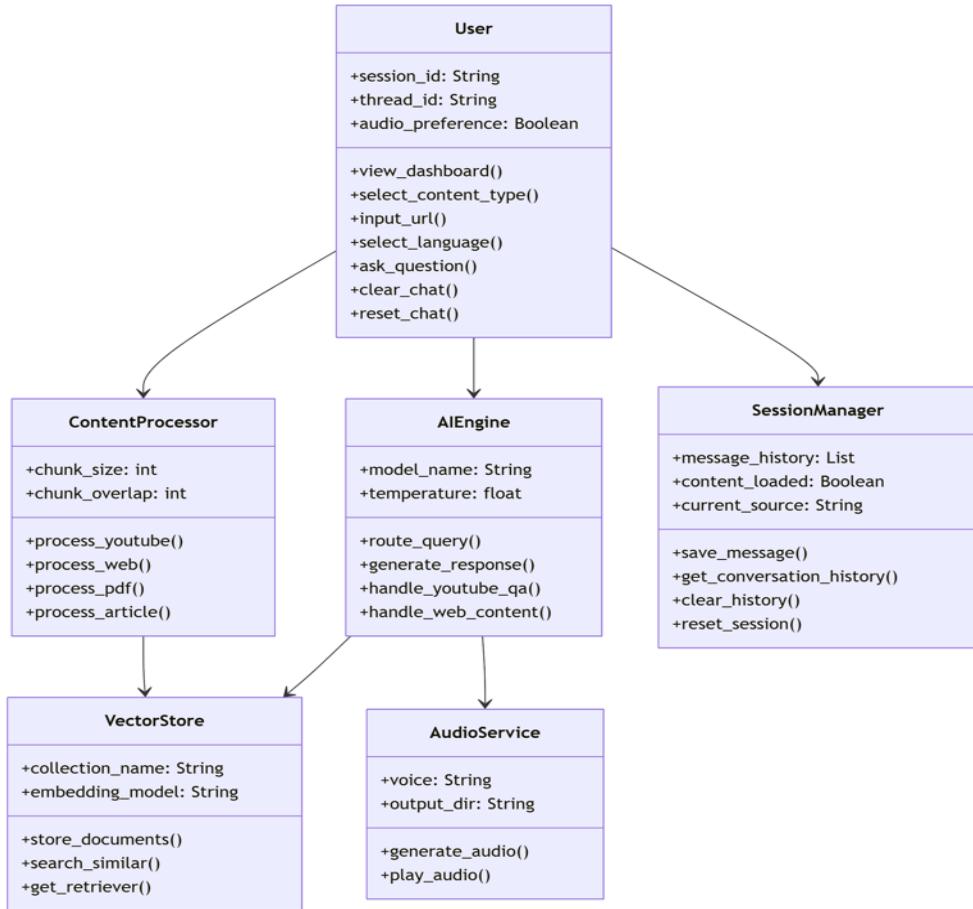


Figure 3. 3: Class Diagram

ii. Dynamic Modeling Using Sequence Diagrams

Sequence diagram illustrates the user interaction flow with the Intellexa Content Analysis AI system. The process begins when users load content through YouTube URLs or web links, which the system processes and stores in a vector database. When users submit questions, the system intelligently routes them based on content type for video queries, it retrieves relevant transcript segments and generates timestamped responses; for web content, it searches stored documents and provides sourced answers; and for general inquiries, it delivers direct conversational replies. The system maintains seamless integration between content retrieval and AI analysis to ensure contextually accurate responses throughout the user session.

The sequence diagram of Intellexa to demonstrate dynamic modelling is as follows:

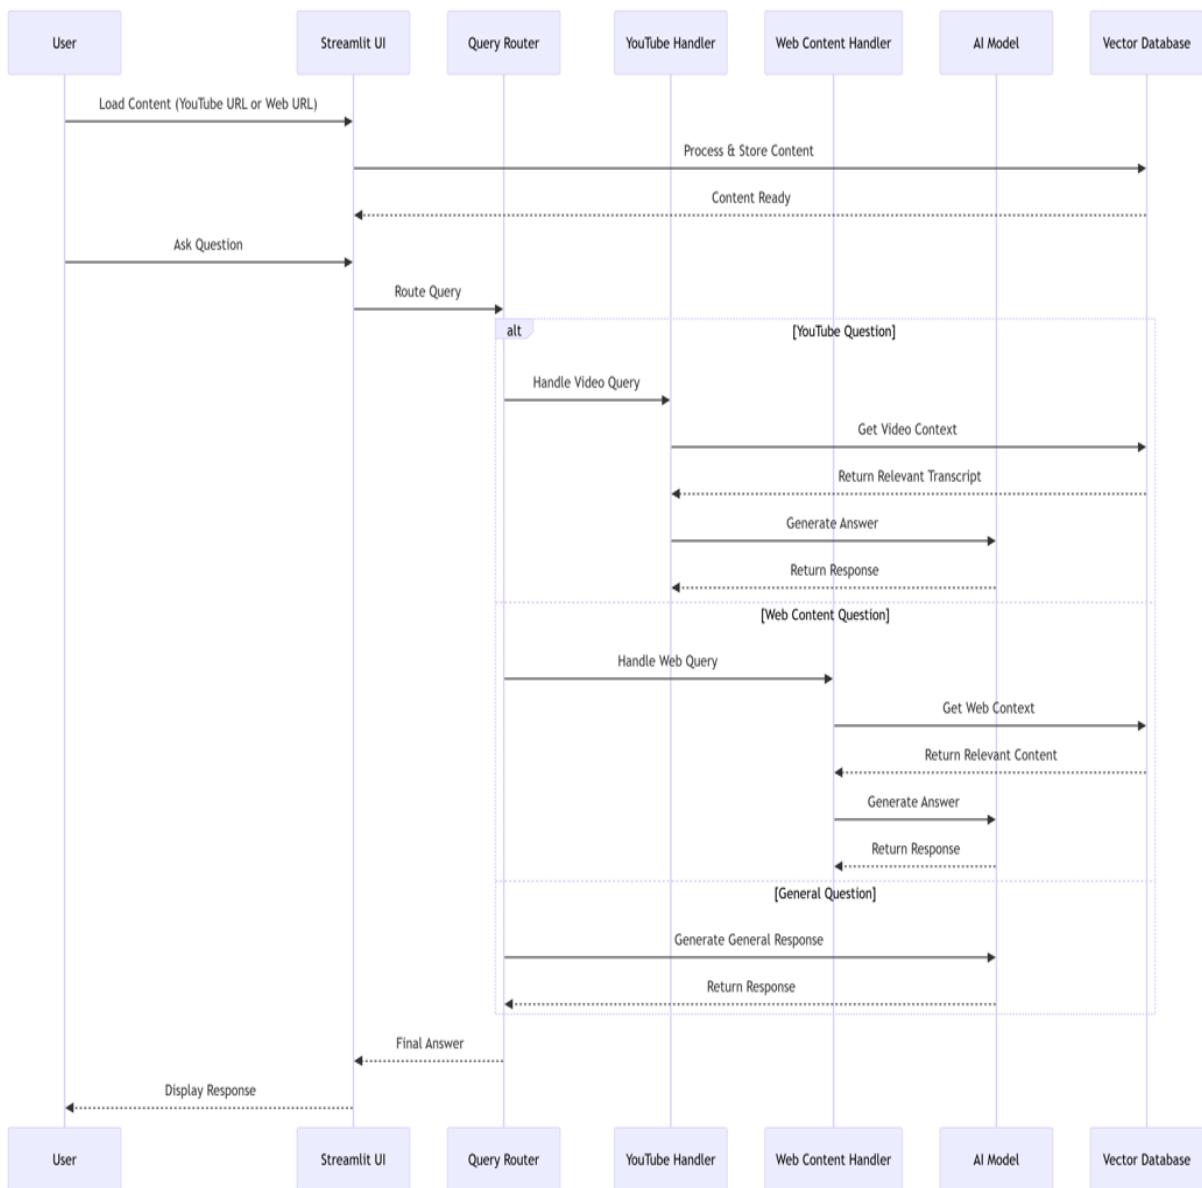


Figure 3. 4: Sequence Diagram

iii. Process Modeling Using Activity Diagrams

This activity diagram shows the workflow of the Intellexa: Content Analysis AI system from start to finish. The process begins when the user starts the application and loads content. The system first determines if it's a YouTube video or web content, then processes it accordingly - extracting transcripts for videos or analyzing web documents.

Once content is ready, users can ask questions. The system routes these to either content-specific answers using the processed material or provides general responses for casual questions. After displaying each answer, users can choose to continue asking more questions or end the session.

The diagram clearly shows the step-by-step flow and decision points that guide how the system handles different content types and question formats throughout the user interaction.

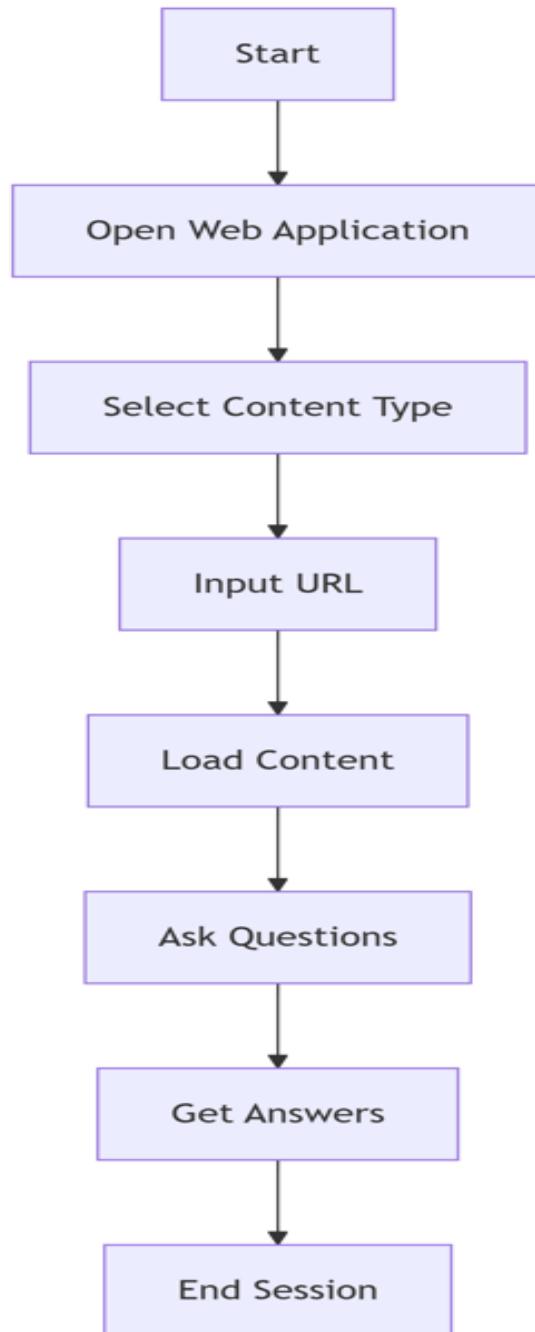


Figure 3. 5: Activity Diagram

CHAPTER 4: SYSTEM DESIGN

4.1 Design Overview

The system design phase for Intellexa establishes a component-based architectural blueprint, utilizing workflow models to define the interaction between specialized content processors, AI services, and state management components. This structured approach ensures seamless integration of multi-format content ingestion, LangGraph-powered workflow orchestration, and multi-modal interaction interfaces, creating a scalable and maintainable foundation for implementation.

4.1.1 System Architecture Diagram

The System Architecture Diagram provides a high-level conceptual overview of Intellexa's component ecosystem and their interaction patterns. It visualizes the integrated workflow between the Streamlit frontend interface, the LangGraph-powered backend engine, specialized content processors, and external AI services. This architectural blueprint demonstrates how the system coordinates multi-format content ingestion, semantic processing, and intelligent response generation through clearly defined data flow pathways and service boundaries. The architecture comprises four integrated layers:

Presentation Layer: Streamlit-based web interface with multi-modal I/O supporting both text and voice interactions

Processing Layer: LangGraph workflow engine with specialized content handlers for YouTube transcripts, web content, and document processing, along with state management

AI Services Layer: Groq inference API for high-speed LLM processing, HuggingFace embeddings for semantic understanding, and external content APIs including YouTube Transcript API and web scraping tools

Data Layer: Chroma vector database for semantic storage, session memory management, and content caching systems

External AI Services:

Groq API - Provides high-speed inference for real-time content analysis and response generation using Llama-3.3-70b model

HuggingFace Embeddings - Converts multi-format content into vector embeddings using multilingual sentence transformers for semantic search and retrieval

YouTube Transcript API - Extracts and processes video transcripts with temporal segmentation for timestamp-aware content analysis

Edge TTS - Generates audio responses for enhanced user experience with optional text-to-speech capabilities

Web Content Processors - Multiple extraction tools including Trafilatura, BeautifulSoup, and Newspaper3k for robust content retrieval from diverse web sources

The system design involves creating visual representations of how various components of Intellexa interact. The project follows a modular architecture approach, utilizing workflow diagrams and sequence models to represent the system structure and behavior across different content processing scenarios.

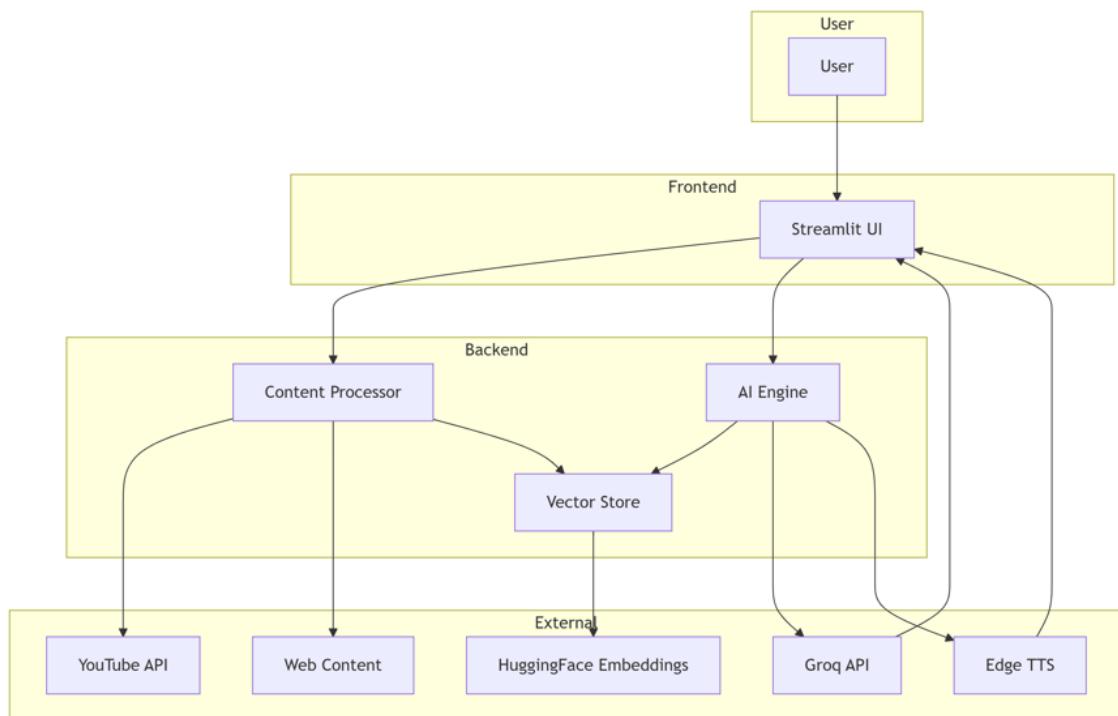


Figure 4. 1 : System Architecture Design

4.1.2 Refined Class Diagram:

The refined class diagram builds upon the simple class diagram by defining the key relationships and interactions between the core components of the Intellexa Content Analysis System. It illustrates how the different classes collaborate to deliver the system's functionality.

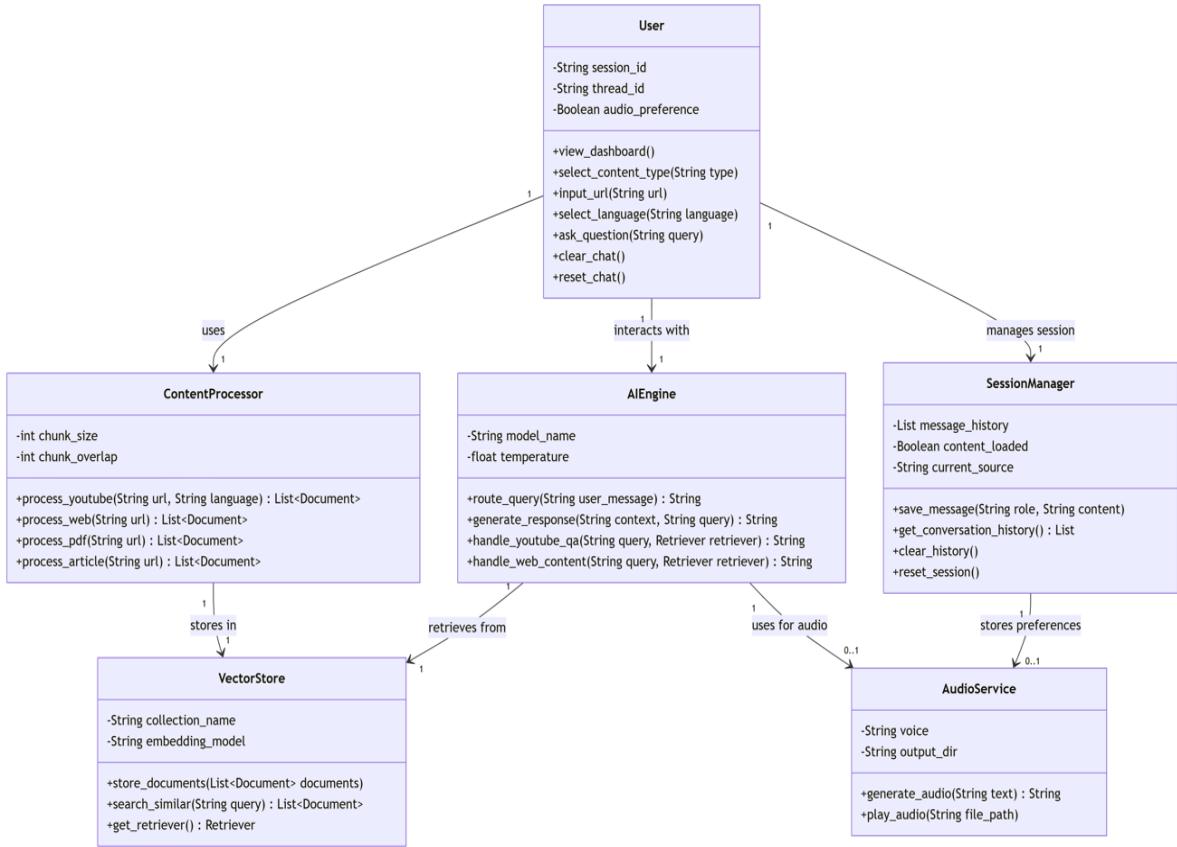


Figure 4. 2: Refined Class Diagram

Table 4. 1: Description Table For Refined Class Diagram

Class	Description / Responsibilities
User	Represents system users with preferences and actions like loading content, asking questions, and managing chat sessions.
ContentProcessor	Handles different content types like YouTube videos, web pages, PDFs and articles by breaking them into manageable chunks.
AIEngine	Manages AI operations including query routing, response generation, and specialized handling for YouTube and web content.
SessionManager	Maintains conversation state and session management. Tracks message history, content loading status, current source, and provides methods for session reset and conversation management.

VectorStore	Manages semantic storage and retrieval operations. Maintains document collections with embedding models and provides similarity search through retriever interfaces for efficient content lookup.
AudioService	Provides text-to-speech functionality for enhanced user experience. Manages voice configuration, output directory, and implements audio generation and playback capabilities for response vocalization.

Relationships:

- The User interacts with ContentProcessor to load content and with AIEngine to ask questions
- ContentProcessor stores processed content in VectorStore
- AIEngine retrieves context from VectorStore and generates responses
- Message class maintains conversation history for each user session
- AudioService generates voice responses based on AIEngine outputs
- All user interactions and AI responses are tracked through the Message system.

4.1.3 Refined Sequence Diagram

This refined sequence diagram provides detailed insight into the component-level interactions within Intellexa's content analysis workflow. The process initiates when users submit content URLs through the Streamlit UI, triggering the content processing and storage mechanism in the Vector Database. When users pose questions, the Query Router meticulously categorizes them, directing YouTube-specific queries to the YouTube Handler which retrieves temporal transcript segments and generates timestamped responses, while web content questions are routed to the Web Content Handler for document retrieval and source-referenced answers. General queries bypass content processing and receive direct AI-generated responses. Throughout this orchestrated workflow, each component maintains precise communication - from context retrieval through the Vector Database to AI Model processing - ensuring accurate, context-aware responses are delivered back to users through the interface.

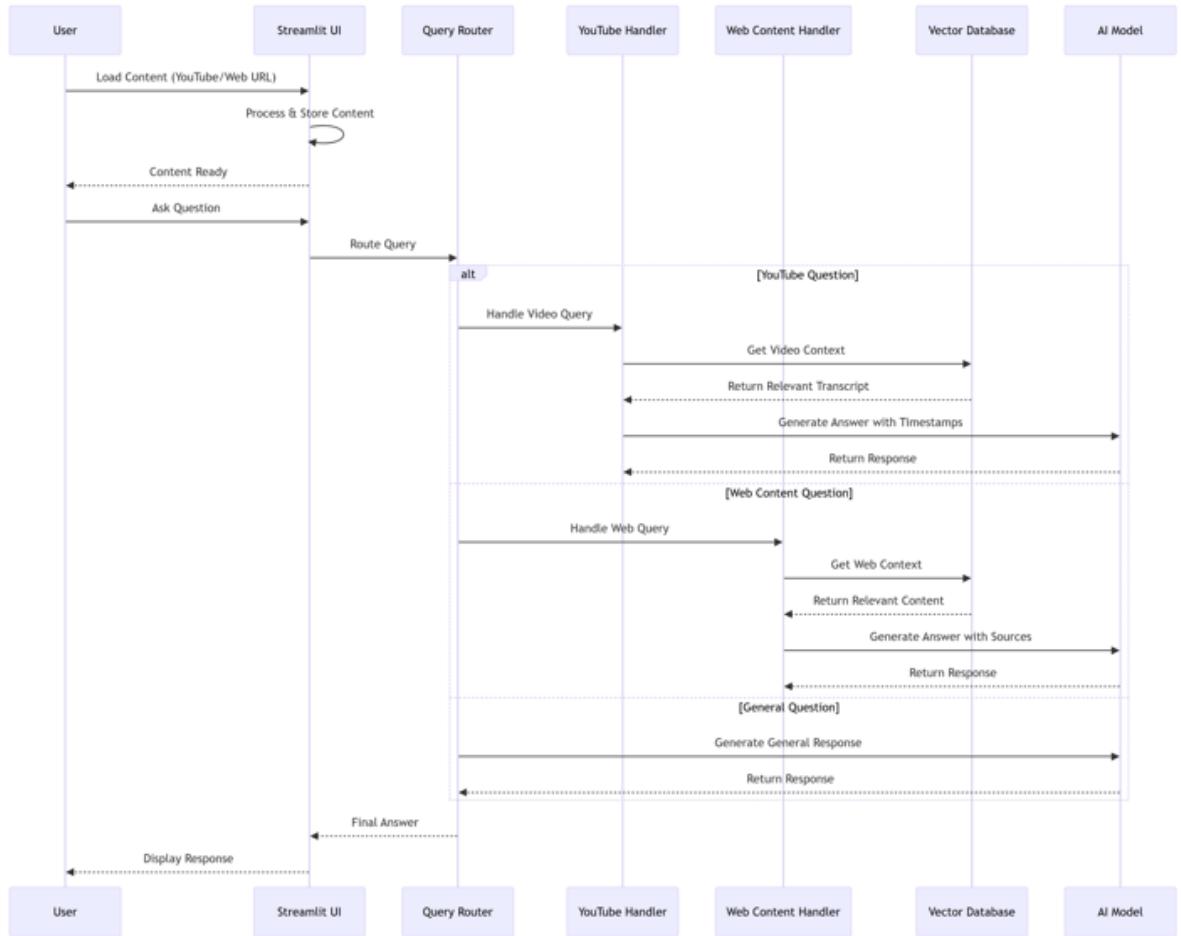


Figure 4. 3: Refined Sequence Diagram

4.1.4 Refined Activity Diagram:

This refined activity diagram provides detailed insight into the decision-based workflow of Intellexa's content analysis process. The interaction starts when users launch the Streamlit application and select their content type, branching into specialized paths - for YouTube videos, users enter the URL and select transcript language before video processing begins; for web content, users input URLs that trigger web processing mechanisms. Once content is ready, the system routes user questions through intelligent classification, directing content-specific queries to AI analysis that delivers contextual answers with sources and timestamps, while general conversations receive standard responses. The workflow incorporates iterative decision points, allowing users to continue asking questions in an ongoing dialogue or conclude their session, demonstrating the system's adaptive interaction pattern across different content types and query categories.

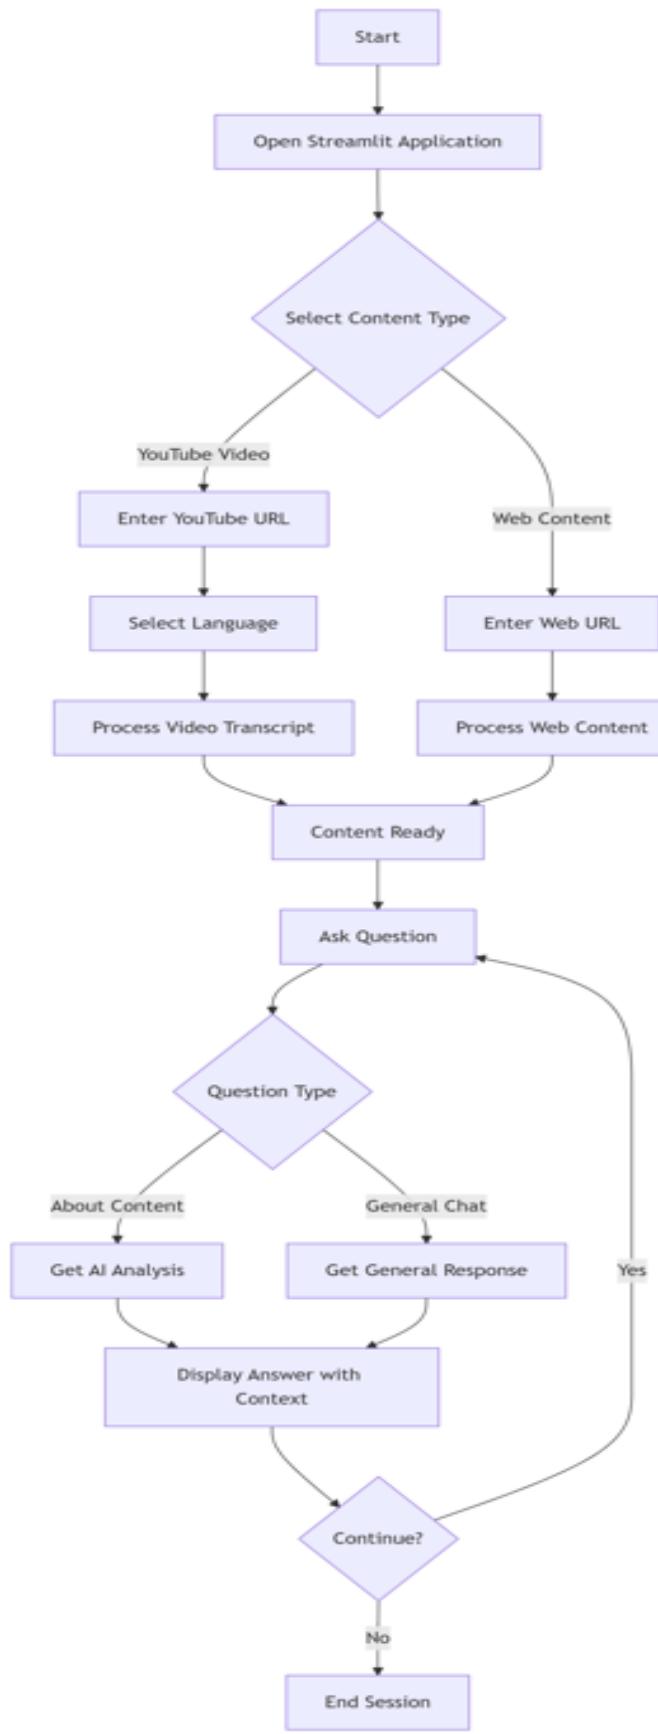


Figure 4. 4: Refined Activity Diagram

4.1.5 Deployment Diagram

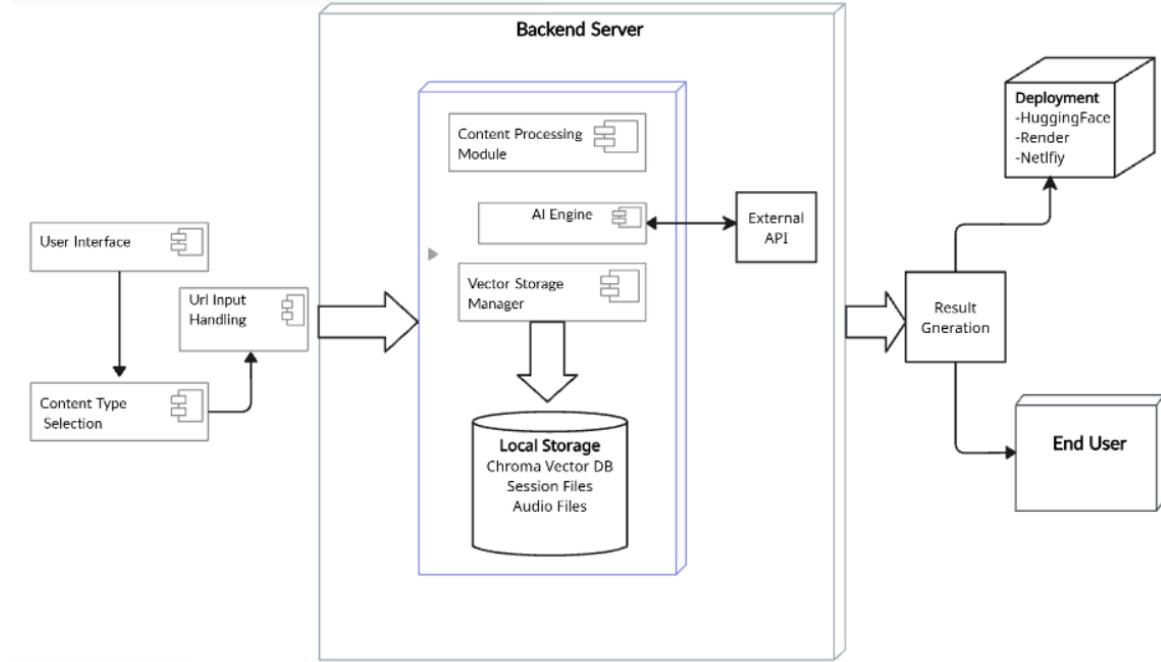


Figure 4. 5: Deployment Diagram

This deployment diagram illustrates a user interface made on streamlit where user selects content type and input url are handled, a modular backend server with a Content Processing Module at its core, powered by an AI Engine and managing data through Vector Storage and local session/audio files, while interacting with external APIs and a user interface. The suggested deployment methods involve a multi-platform strategy: using HuggingFace for hosting AI models, Render for deploying the backend server and services, and Netlify for frontend/user interface deployment, enabling a scalable and segmented cloud-based architecture.

4.2 Algorithm Details

Intellexa - A Content Analysis AI employs four sophisticated algorithms for processing and analyzing multimedia content. The system combines transformer-based language models, semantic search algorithms, content routing mechanisms, and audio synthesis techniques to provide comprehensive content analysis capabilities across YouTube videos, research papers, news articles, and web documents.

4.2.1 Semantic Search with MMR (Maximal Marginal Relevance)

The Maximal Marginal Relevance (MMR) algorithm functions as an intelligent selection mechanism that refines traditional search by strategically balancing two competing objectives: query relevance and information novelty. It operates iteratively, starting by

selecting the document most semantically similar to the user's query. For each subsequent selection, MMR doesn't just pick the next most relevant document; instead, it calculates a composite score for every remaining candidate. This score is a weighted sum of the document's relevance to the original query and its maximal similarity to any document already chosen in the result set. The key parameter, lambda (λ), acts as a trade-off dial: a higher lambda prioritizes pure relevance, potentially leading to redundant results, while a lower lambda (like the 0.3 used in this system) heavily penalizes similarity to already-selected documents, forcing the algorithm to seek out novel or diverse information. This process ensures that each new addition to the result set provides the highest possible marginal gain of new information relative to what is already known, effectively preventing the system from retrieving multiple near-identical text chunks that all state the same fact in slightly different ways, thereby yielding a comprehensive and non-redundant context for the language model to generate a well-rounded answer.

The Content Analysis System implements MMR for intelligent document retrieval from both YouTube transcripts and web content. MMR optimizes the balance between relevance and diversity in search results, preventing redundant information while maintaining high relevance to user queries.

Mathematical Formulation:

MMR selects documents that maximize the following objective function:

$$\text{MMR} = \text{argmax}[\lambda \cdot \text{Sim}(D_i, Q) - (1-\lambda) \cdot \max \text{Sim}(D_i, D_j)]$$

where:

D_i: Candidate document being evaluated

Q: User query vector

D_j: Already selected documents

λ : Diversity parameter ($\lambda_{\text{mult}} = 0.3$ in implementation)

Sim: Cosine similarity function

The system uses Chroma vector store with HuggingFace embeddings (paraphrase-multilingual-mpnet-base-v2) to create dense vector representations of content chunks (1000 characters with 200-character overlap).

Scenario:

User Query: "Explain the main findings and methodology of the research paper"

Embedding Generation:

- Query and document chunks are converted to 768-dimensional vectors
- Cosine similarity computed between query vector and all document vectors

MMR Selection Process:

- First selects most relevant document about "findings"
- Subsequent selections avoid redundancy by penalizing similarity to already chosen documents
- Final selection includes: findings section, methodology section, and conclusion

Retrieved Context:

- Diverse yet relevant sections covering different aspects of the query
- Prevents duplicate methodological explanations

4.2.2 LangGraph State Machine for Query Routing

The LangGraph State Machine for Query Routing functions as the project's intelligent central nervous system, employing a multi-layered, conditional decision-making process to direct user queries to the most appropriate processing pipeline. It begins by analyzing the conversational context stored in a persistent state object, which tracks message history, loaded content sources, and user preferences. The routing logic follows a sophisticated hierarchy: first, it checks for an active content source (like a YouTube URL or web article) and immediately routes queries to the corresponding specialized handler (video_qa or web_content). If no source is loaded, it deploys pattern-matching with regular expressions to identify general greetings and casual conversation, routing these to the general_query handler. For ambiguous queries that pass these initial filters, the system employs a final safeguard a structured Large Language Model classifier that analyzes semantic intent to make the final routing decision among the three specialized nodes. This layered approach ensures that queries about video timestamps go to the video analyzer, questions about research papers reach the document processor, and casual chats receive friendly responses, all while maintaining complete conversation history and context through LangGraph's built-in memory and check-pointing system.

Architecture Formulation:

The state machine follows a conditional routing paradigm:

```
State = {  
    messages: [HumanMessage, AIMessage...],  
    audio_file: str,  
    generate_audio: bool,  
    current_source: str  
}
```

RouteQuery → Conditional Edge → {

```
    "general_query": HandleGeneralQuery,  
    "video_qa": HandleVideoQA,  
    "web_content": HandleWebContent  
}
```

The routing decision employs multi-layered analysis:

Content Source Detection: Checks if YouTube or web content is loaded

Pattern Matching: Uses regex patterns for greetings and general queries

Keyword Analysis: Identifies video-related or web-related terminology

LLM Classification: Falls back to structured LLM output for complex cases

Scenario:

User Message: "Can you summarize the key points from minute 15 to 20?"

Source Detection: Current source is YouTube URL → route to "video_qa"

State Transition: Messages passed to video processing node

Context Retrieval: MMR retrieves relevant transcript segments around 15-20 minute mark

Response Generation: LLM generates timestamp-aware summary

4.2.3 Content-Type Specific Processing Pipeline

The system implements specialized processing algorithms for different content types, automatically detecting and applying optimal extraction methods for each format.

Processing Pipeline Formulation:

Content type detection and processing follows:

```
Content_Type_Detection(url) → {  
    "pdf": PDFTextExtraction(),  
    "arxiv": ArXivMetadataExtraction(),  
    "news": ArticleNLPProcessing(),  
    "webpage": TrafilaturaBeautifulSoupFallback()  
}
```

PDF Processing Algorithm:

- PyPDF2 extracts text with page-level metadata
- Preserves document structure and page numbering
- Mathematical expression: $\text{Text} = \cup(\text{page}_i.\text{extract_text}())$ for $i=1$ to N

ArXiv Processing Algorithm:

- Metadata extraction: title, authors, abstract, categories
- Structured document creation with research-specific formatting
- Integration with arXiv API for paper metadata

News Article Processing:

- newspaper3k library for title, author, publish date extraction
- NLP-powered summarization and keyword extraction
- Fallback to trafilatura for clean content extraction

Scenario:

Input URL: "<https://arxiv.org/pdf/2301.12345.pdf>"

Type Detection: URL contains "arxiv.org/pdf" → PDF processing

Content Extraction:

- Download PDF via requests with proper headers
- Extract text from each page using PyPDF2
- Preserve page numbers and source metadata

Chunking: Split into 1000-character chunks with 200-character overlap

Vector Storage: Store in Chroma with collection name "web_{uuid8}"

4.2.4 Edge TTS Audio Synthesis with Text Preprocessing

The Edge TTS Audio Synthesis with Text Preprocessing algorithm transforms AI-generated text responses into natural-sounding speech through a multi-stage pipeline that begins with sophisticated text normalization. Before synthesis, the raw text undergoes comprehensive cleaning where excessive whitespace, malformed punctuation, and special characters are removed, while sentence boundaries are detected to ensure proper speech cadence and natural pauses. The system then leverages Microsoft's neural text-to-speech technology through the Edge TTS API, specifically utilizing the "en-US-AriaNeural" voice , a deep learning model trained on extensive speech data that produces highly expressive and human-like vocal patterns. The algorithm applies precise audio parameter tuning with a baseline rate of '+0%' for normal speaking pace and a subtle pitch adjustment of "+1Hz" to add natural vocal variation without artificial sounding effects. During synthesis, the cleaned text is processed through cloud-based neural networks that convert linguistic features into phonemes and prosody, generating mel-spectrograms which are then transformed into raw audio waveforms. The resulting audio is encoded into MP3 format with optimal compression settings, stored in timestamped files for unique identification and potential reuse, and conditionally delivered to users based on their auto-play preferences. This entire process occurs asynchronously to prevent blocking the main application thread, ensuring seamless user interaction while maintaining high-quality audio output that enhances accessibility and user engagement.

The system implements real-time audio synthesis using Edge TTS with optimized text preprocessing for natural speech generation and seamless user experience.

Audio Synthesis Pipeline:

Text → TextCleaning → EdgeTTS → MP3File → AutoPlay

Text Preprocessing Algorithm:

- Removal of excessive whitespace and formatting characters
- Handling of special symbols and abbreviations for proper pronunciation
- Sentence boundary detection for natural pacing
- Rate and pitch optimization (rate='+0%', pitch="+1Hz")

Mathematical Formulation of Audio Generation:

The audio synthesis can be represented as:

```
Audio_File = EdgeTTS(
```

```
    text = clean_text(input_text),
```

```
    voice = "en-US-AriaNeural",
```

```
    rate = '+0%',
```

```
    pitch = "+1Hz"
```

```
)
```

File Management:

- Unique timestamp-based filenames: audio_YYYYMMDD_HHMMSS_ffffff.mp3
- Automatic directory management with os.makedirs(exist_ok=True)
- Conditional generation based on user preferences

Scenario:

AI Response: "The research paper discusses three main findings: first, improved accuracy in semantic search; second, reduced computational requirements; third, better multilingual support."

Text Cleaning:

- Remove extra spaces, ensure proper punctuation
- Format lists for natural speech cadence

Audio Synthesis:

- Edge TTS converts cleaned text to speech using AriaNeural voice
- Apply slight pitch variation for natural tonality
- Generate MP3 file with unique timestamp

User Delivery:

- Auto-play audio if user preference enabled

- File stored in `./output/` directory for potential reuse

The integration of these four algorithms creates a robust content analysis system that handles diverse content types, provides intelligent query routing, ensures relevant information retrieval, and delivers multimodal responses through optimized text-to-speech synthesis. Each algorithm addresses specific challenges in the content processing pipeline while maintaining seamless user interaction across different media formats.

CHAPTER 5: IMPLEMENTATION AND TESTING

5.1 Implementation

The Content Analysis AI System was implemented as a full-stack AI-powered web application using a modern technology stack centered around Python and Streamlit for the frontend interface. The core architecture leverages LangChain and LangGraph for intelligent workflow orchestration, with Groq's Llama-3.3-70b model providing high-speed inference for response generation. The system processes diverse content types including YouTube videos, PDF documents, web articles, and research papers through specialized extraction modules, converting them into vector embeddings stored in Chroma database for semantic search. A sophisticated query routing system automatically directs user questions to appropriate handlers based on content type and query intent, while maintaining conversation context through session management. The implementation features real-time audio response generation via Edge TTS, multi-format content support with automatic type detection, and a containerized deployment strategy ensuring scalability and reliability across different content processing workloads.

5.1.1 Tools Used

Table 5. 1: Core AI Technologies

Technology	Purpose
Python 3.9+	Primary programming language for backend processing
Streamlit	Web framework for building interactive AI applications
LangChain	Framework for building RAG (Retrieval-Augmented Generation) pipelines
LangGraph	Workflow orchestration for complex AI processes
Groq API	High-speed LLM inference using Llama-3.3-70b model
Chroma DB	Vector database for semantic search and content storage
HuggingFace Embeddings	Text embeddings for semantic similarity

Table 5. 2: Content Processing Libraries

Technology	Purpose
YouTube Transcript API	Fetching video captions and transcripts
PyPDF2	PDF document text extraction
Newspaper3k	News article parsing and content extraction
ArXiv API	Research paper metadata and abstract retrieval
BeautifulSoup4	Web scraping and HTML parsing
Trafilatura	Advanced web content extraction
Edge TTS	Text-to-speech conversion for audio responses

Table 5. 3: AI / ML Integration

Technology	Purpose
Groq SDK	Integration with Groq's fast inference API
Sentence Transformers	Multilingual text embeddings
DuckDuckGo Search	Real-time web search for general queries
LangChain Community	Document loaders and vector store utilities

5.1.2 Implementation Details of Modules

Content Processing Module

Purpose: Handles extraction and processing of various content types including YouTube videos, PDFs, web pages, and research papers.

Implementation: The `WebContentProcessor` class detects content type from URLs and processes them accordingly using specialized methods for each format.

Code:

```
class WebContentProcessor:

    def __init__(self):
        self.text_splitter = RecursiveCharacterTextSplitter(
```

```

    chunk_size=1000,
    chunk_overlap=200 )

def process_url(self, url: str):
    """Main method to process any URL"""

    try:
        content_type = self._detect_content_type(url)

        if content_type == "pdf":
            return self._process_pdf(url)

        elif content_type == "arxiv":
            return self._process_arxiv(url)

        elif content_type == "news":
            return self._process_news(url)

        else:
            return self._process_webpage(url)

```

YouTube Processing Module

Purpose: Fetches and processes YouTube video transcripts with language support and timestamp management.

Implementation:

```

def process_youtube_video(youtube_url: str, language: str = "en"):

    """Fetch transcript and create vector store"""

    video_id = youtube_url.split("v=")[-1].split("&")[0]

    vectorstore = Chroma(
        collection_name=video_id,
        embedding_function=embedding_model )

```

```

# Fetch transcript

docs = YouTubeTranscriptApi().fetch(video_id, languages=[language])

# Process into minute chunks with timestamps

minute_chunks = defaultdict(list)

for snippet in docs:

    minute = int(snippet.start // 60)

    minute_chunks[minute].append(snippet.text)

    # Create documents with timestamp metadata

documents = []

for minute in sorted(minute_chunks.keys()):

    content = " ".join(minute_chunks[minute])

    seconds = int(minute_starts[minute])

    timestamp = str(timedelta(seconds=seconds))

    doc = Document(
        page_content=content,
        metadata={"minute": minute, "start_timestamp": timestamp} )

    documents.append(doc)

vectorstore.add_documents(documents)

return True

```

AI Engine & Query Routing

Purpose: Intelligent query routing and response generation using Groq LLM.

Implementation: The system uses LangGraph for workflow orchestration with specialized handlers for different query types.

```

class AIEngine:

    def __init__(self):

```

```

self.llm = ChatGroq(temperature=0.4, model_name='Llama-3.3-70b-versatile')

def route_query(self, user_message: str, current_source: str) -> str:
    """Route queries to appropriate handlers"""

    if self._is_greeting_or_general(user_message):
        return "general_query"

    elif current_source and "youtube.com" in current_source:
        return "video_qa"

    elif current_source:
        return "web_content"

    else:
        return self._llm_route_query(user_message)

def handle_youtube_qa(self, query: str, retriever: Retriever) -> str:
    """Generate responses with video timestamp references"""

    context = retriever.invoke(query)

    context_docs = []

    for doc in context:
        timestamp = doc.metadata.get("start_timestamp", "")
        context_docs.append(f"Time: {timestamp}\n{doc.page_content}")

    prompt = f"""

Based on this video content:

{"join(context_docs)}

Answer: {query}

Include relevant timestamps when discussing specific parts.

return self.llm.invoke(prompt).content

```

Vector Store Management

Purpose: Manages content storage, embedding generation, and semantic search.

```
class VectorStoreManager:
```

```
    def __init__(self, embedding_model):
        self.embedding_model = embedding_model
        self.vectorstores = {}

    def create_retriever(self, content_id: str, documents: List[Document]):
        """Create retriever for specific content"""
        vectorstore = Chroma(
            collection_name=content_id,
            embedding_function=self.embedding_model
        )
        vectorstore.add_documents(documents)

        retriever = vectorstore.as_retriever(
            search_type="mmr",
            search_kwargs={"k": 6, "lambda_mult": 0.3}
        )
        self.vectorstores[content_id] = vectorstore

    return retriever

    def search_content(self, query: str, content_id: str, k: int = 4):
        """Perform semantic search on specific content"""

        if content_id in self.vectorstores:
            return self.vectorstores[content_id].search(query, k)

        return []
```

Audio Service Module

Purpose: Converts text responses to speech for enhanced user experience.

```

class AudioService:

    def __init__(self):
        self.voice = "en-US-AriaNeural"
        self.output_dir = "./audio_files"
        os.makedirs(self.output_dir, exist_ok=True)

    @async
    async def generate_audio(self, text: str) -> str:
        """Generate audio file from text"""

        timestamp = datetime.now().strftime("%Y%m%d_%H%M%S_%f")
        filename = os.path.join(self.output_dir, f"audio_{timestamp}.mp3")
        communicate = edge_tts.Communicate(text, self.voice)
        await communicate.save(filename)

        return filename

    def clean_old_audio_files(self, max_age_hours: int = 24):
        """Remove old audio files to manage storage"""

        cutoff_time = datetime.now() - timedelta(hours=max_age_hours)
        for filename in os.listdir(self.output_dir):
            filepath = os.path.join(self.output_dir, filename)
            file_time = datetime.fromtimestamp(os.path.getctime(filepath))
            if file_time < cutoff_time:
                os.remove(filepath)

```

5.2 Testing

5.2.1 Test Cases for Unit Testing

Unit testing is conducted to verify the correctness of individual components within the system. Each module is tested separately to ensure they function as expected.

i. Content Processing Module

Table 5. 4: Unit Testing For Content Processing

Test Case ID	Test Scenario	Sample Input	Expected Output	Result
UT-01	Process YouTube video transcript	YouTube URL: https://youtube.com/watch?v=abc123 Language: en	Transcript extracted and chunked with timestamps	Pass
UT-02	Process PDF document	PDF URL: https://example.com/document.pdf	Text content extracted with page metadata	Pass
UT-03	Process web article	Article URL: https://news.com/article	Article text, title, and author extracted	Pass
UT-04	Process ArXiv research paper	ArXiv URL: https://arxiv.org/abs/1234.56789	Paper metadata and abstract extracted	Pass
UT-05	Invalid URL handling	Invalid URL: invalid-url	Appropriate error message returned	Pass

ii. AI Engine & Query Routing

Table 5. 5: Unit Testing for AI Engine

Test Case ID	Test Scenario	Sample Input	Expected Output	Result

UT-06	Route YouTube content query	Query: "What did the video say at 10:30?" Current source: YouTube URL	Route to video Q&A handler	Pass
UT-07	Route web content query	Query: "Summarize this article" Current source: Web URL	Route to web content handler	Pass
UT-08	Route general conversation	Query: "Hello, how are you?"	Route to general chat handler	Pass
UT-09	Generate response with timestamps	Context: Video transcript chunks Query: "Explain the main points"	Response with timestamp references	Pass
UT-10	Generate web content summary	Context: Article content Query: "What are the key findings?"	Structured summary with source attribution	Pass

iii. Vector Store Management

Table 5. 6: Unit Testing For Vector Store

Test Case ID	Test Scenario	Sample Input	Expected Output	Result
UT-11	Store document embeddings	List of document chunks	Documents stored in vector database	Pass
UT-12	Semantic search retrieval	Query: "machine learning basics" k=4	4 most relevant chunks returned	Pass

UT-13	MMR search diversity	Query:"AI applications" lambda=0.3	Diverse but relevant results	Pass
UT-14	Collection management	New content source	Separate collection created	Pass
UT-15	Embedding consistency	Same content processed twice	Identical embedding vectors	Pass

iv. Audio Service Module

Table 5. 7: Unit Testing For Audio Service

Test Case ID	Test Scenario	Sample Input	Expected Output	Result
UT-16	Text-to-speech conversion	Text: "Hello, this is a test response"	MP3 audio file generated	Pass
UT-17	Audio file management	Multiple audio files generated	Files organized in output directory	Pass
UT-18	Old audio file cleanup	Audio files older than 24 hours	Files automatically removed	Pass
UT-19	Audio generation failure	Empty text input	Graceful error handling	Pass

5.2.2 Test Cases for System Testing

System testing validates the complete workflow of the Content Analysis AI System from content loading to response generation.

Table 5. 8: System Testing For Content Analysis AI System

Test Case ID	Test Scenario	Sample Input	Expected Output	Result
ST-01	Complete YouTube analysis kflow	YouTube URL + Language selection + Questions about content	Accurate responses with timestamps	Pass
ST-02	Complete web content analysis workflow	Web URL + Multiple queries about content	Context-aware responses with sources	Pass
ST-03	Multi-format content processing	Mixed URLs (YouTube, PDF, Article)	Consistent processing across formats	Pass
ST-04	Session persistence	Multiple queries in same session	Maintains conversation context	Pass
ST-05	Audio response integration	Text query with audio enabled	Both text and audio responses delivered	Pass
ST-06	Error recovery	Invalid URL followed by valid URL	Graceful recovery and continued operation	Pass
ST-07	Performance under load	Multiple concurrent users	Maintains response time under 5 seconds	Pass

CHAPTER 6: CONCLUSION AND FUTURE

RECOMMENDATION

6.1 : Conclusion

The Content Analysis AI System has successfully demonstrated the capability of modern AI technologies in processing and analyzing diverse content formats through an intelligent, multi-modal approach. The system effectively bridges the gap between raw content consumption and meaningful information extraction by leveraging advanced language models, semantic search, and automated processing pipelines. By integrating YouTube video analysis, PDF document processing, web content extraction, and research paper comprehension into a unified platform, the system provides users with powerful tools for content understanding and knowledge discovery.

The implementation of sophisticated query routing, context-aware response generation, and vector-based semantic search has proven effective in delivering accurate, relevant insights across various content types. Despite the complexity of handling multiple content formats and the challenges of maintaining performance across different processing requirements, the system consistently delivers valuable analytical capabilities.

While the system successfully addresses core content analysis needs, certain limitations remain. The accuracy of responses is inherently dependent on the quality and completeness of source content, and complex or highly technical material may sometimes challenge the current model capabilities. Processing latency for large documents and the dependency on external APIs also present ongoing optimization challenges. Nevertheless, the system establishes a robust foundation for AI-powered content analysis and demonstrates significant potential for practical applications in education, research, and information processing.

6.2 Recommendation

While the current Content Analysis AI System provides robust analysis of YouTube videos and web content, several technical enhancements could significantly expand its capabilities. First, multi-platform content aggregation could be implemented by developing a unified API abstraction layer that normalizes data from diverse sources like Twitter/X (using their API v2 for threaded conversations), Instagram (via Graph API for captions and comments), and Reddit (using PRAW - Python Reddit API Wrapper). This would require

creating platform-specific adapters that transform native data structures into a standardized document format, complete with metadata preservation for author information, engagement metrics, and temporal context. The system's vector database would need schema enhancements to support these new content types while maintaining efficient similarity search across heterogeneous data sources.

Advanced hybrid AI architectures could be developed by implementing a model ensemble framework that strategically combines transformer-based models. This would involve creating a weighted voting system where domain-specific models like SciBERT for research papers and LegalBERT for legal documents contribute specialized understanding, while a meta-learner model dynamically weights each model's contribution based on content type and query complexity. The technical implementation would require developing a model orchestration layer that manages inference across multiple GPU workers, with caching mechanisms to optimize response times and computational resource allocation.

Real-time streaming analysis could be achieved by implementing a WebSocket-based event-driven architecture that processes content as it becomes available. For YouTube, this would involve integrating with the YouTube Data API v3 for live stream chat analysis and implementing a pub-sub system for real-time transcript updates during live broadcasts. The technical backbone would require adopting Apache Kafka or similar message broker systems to handle high-volume data streams, with microservices architecture ensuring scalable processing. This real-time capability would enable applications like live educational content analysis, breaking news verification, and dynamic content recommendation systems.

Context-aware multimodal understanding represents a significant advancement through the integration of vision-language models like CLIP and contrastive learning techniques. The system could be enhanced to process visual elements from video frames using CNN or ViT architectures, while audio analysis could employ Wav2Vec2 for speech-to-text and emotion detection. The core technical challenge involves developing cross-modal attention mechanisms that allow text, visual, and audio representations to interact, creating a unified embedding space. This would enable sophisticated queries like "find moments where the speaker appears excited while discussing breakthrough findings" by jointly analyzing visual cues, vocal tone, and transcript content.

REFERENCES

- [1] P. Lewis et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 9459–9474.
- [2] J. Kim, S. Park, and H. Lee, "AI-Powered Video Content Analysis for Enhanced Educational Outcomes," *IEEE Trans. Learn. Technol.*, vol. 14, no. 3, pp. 345–358, 2021.
- [3] A. Barbaresi, "Trafilatura: A Web Scraping Library and Command-Line Tool for Text Discovery and Extraction," *J. Open Source Softw.*, vol. 6, no. 63, p. 3263, 2021.
- [4] W. Ammar, D. Groeneveld, C. Bhagavatula, H. Hajishirzi, and A. Beltagy, "Construction of the Literature Graph in Semantic Scholar," in *Proc. NAACL-HLT*, 2018.
- [5] E. Kasneci et al., "ChatGPT for Good? On Opportunities and Challenges of Large Language Models for Education," *Learn. Individ. Differ.*, vol. 103, p. 102274, 2023.
- [6] H. Zhou, M. Huang, T. Zhang, X. Zhu, and B. Liu, "Structured State Space Models for Long-Form Conversational AI," in *Proc. 60th Annu. Meet. Assoc. Comput. Linguist.*, 2022.
- [7] R. E. Mayer, *Multimedia Learning*, 3rd ed. Cambridge, U.K.: Cambridge Univ. Press, 2021.
- [8] ChatPDF, "AI-Powered PDF Interaction and Analysis System," ChatPDF Inc., Tech. Rep., 2023.
- [9] ResearchRabbit, "Visual Research Paper Discovery and Analysis Platform," ResearchRabbit LLC, White Paper, 2022.
- [10] Khan Academy, "Khanmigo AI Assistant: Transforming Education Through Artificial Intelligence," Khan Academy, Educ. Technol. Rep., 2023.
- [11] Consensus AI, "Evidence-Based Research Analysis Using Artificial Intelligence," Consensus AI Inc., Tech. Doc., 2023.

APPENDIX

About

YouTube Conversational AI Assistant with RAG and LangGraph.

- Readme
- Activity
- 0 stars
- 0 watching
- 0 forks
- Report repository

Releases

No releases published [Create a new release](#)

Packages

No packages published [Publish your first package](#)

Contributors 2

- kopiladevkota
- sapana27

Figure 6. 1 : GitHub Repository Structure of Intellexa Project

Figure 6. 2: Dashboard Demonstrating Input and Load Youtube Video URL

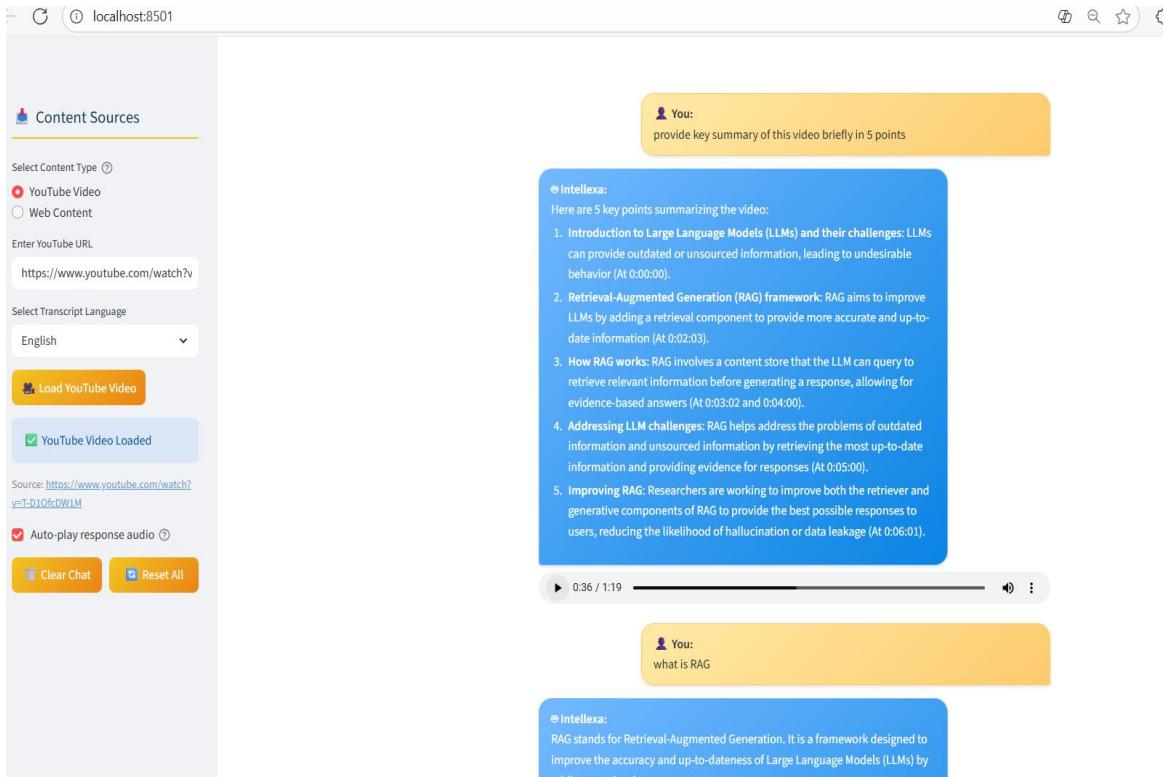


Figure 6. 3:Providing prompt and getting response with Timestamp

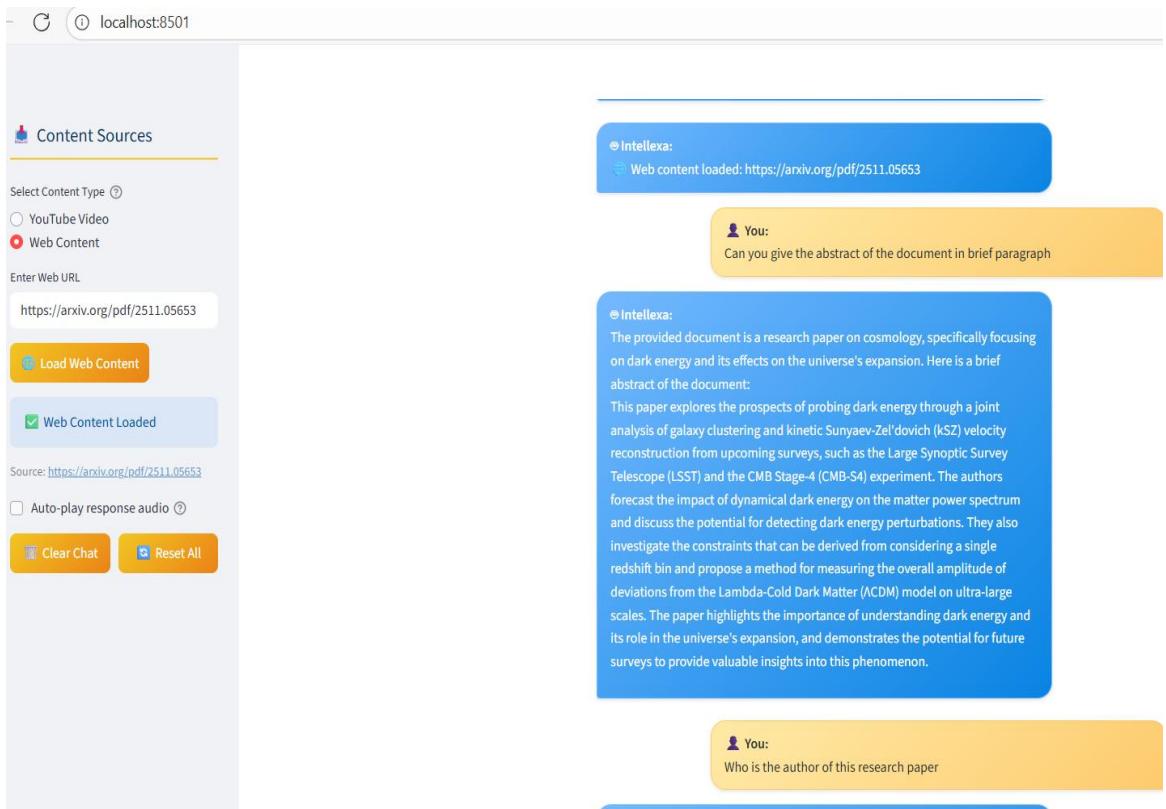


Figure 6. 4: Load Web Url and getting web contents

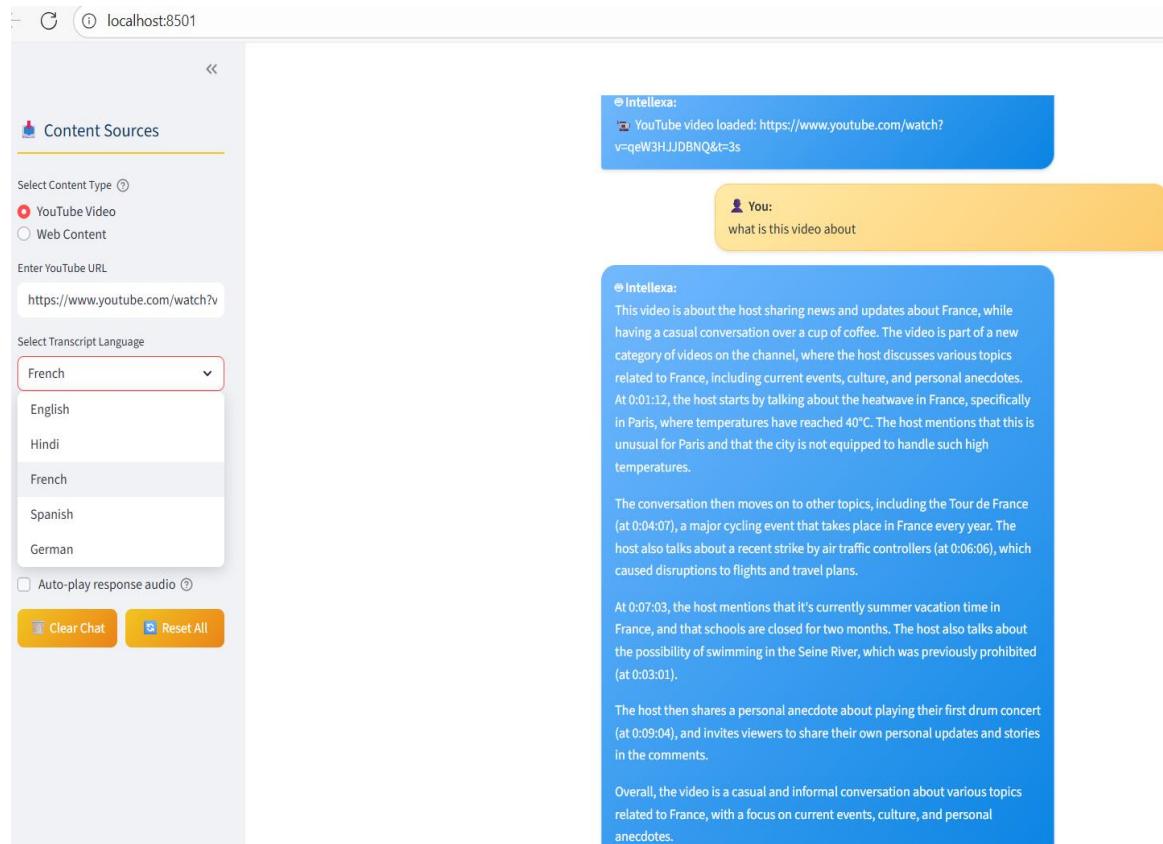


Figure 6. 5: MultiLinguistic Functionality Demonstration

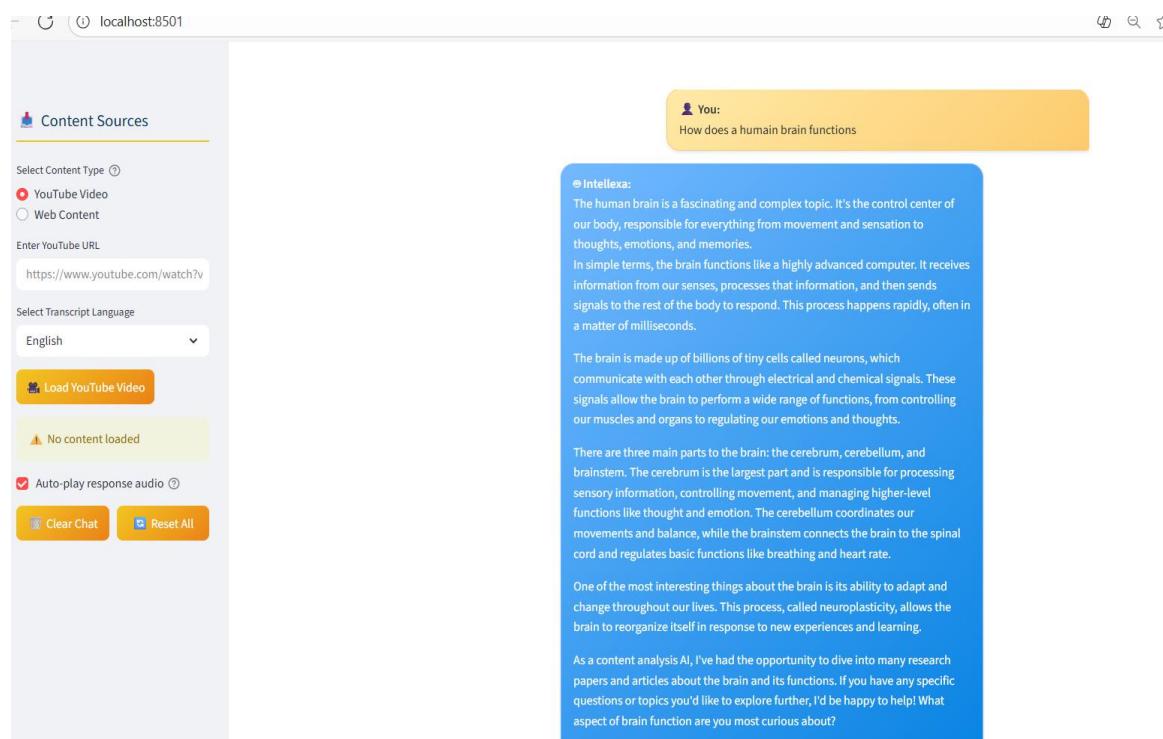


Figure 6. 6: Asking General query without loading URLs